

Generative AI: Current Trends and Practical Applications

Kao Panboonyuen
panboonyuen.kao@gmail.com

<https://kaopanboonyuen.github.io>

Reference:

1. <https://pytorch.org/tutorials>
2. <https://stanford.edu/~shervine/teaching/cs-229/>
3. <http://introtodeeplearning.com/>
4. <https://www.simplilearn.com/tutorials/deep-learning-tutorial/introduction-to-deep-learning>
5. <https://www.geeksforgeeks.org/introduction-deep-learning/>
6. <https://www.telecats.com/blog-en/ai-for-rookies/>
7. <https://marketoonist.com/2023/01/ai-tidal-wave.html>
8. <https://www.v7labs.com/blog/self-supervised-learning-guide>
9. <https://blog.miarec.com/contact-centers-ai-definition>
10. <https://www.solulab.com/top-generative-ai-trends/>

Outlines

- About Me
- Introduction to Generative AI
- Key Trends in Generative AI
- Applications of Generative AI

About Me



Kao
Panboonyuen

kao-panboonyuen 

AI Research Scientist

Name: Teerapong Panboonyuen (P'Kao)

Contact: teerapong.pa@chula.ac.th

panboonyuen.kao@gmail.com

Education: Ph.D. in Computer Engineering, Chula

Position: AI Team Lead, MARS

PostDoc, Chula

Visiting Lecturer, KKU

Interests: Computer Vision, Deep Learning

Machine Learning, Remote Sensing



Teerapong Panboonyuen

Senior Research Scientist at MARS, Post-doc at Chula
Verified email at chula.ac.th - [Homepage](#)


Artificial Intelligence Machine Learning Deep Learning Computer Vision Remote Sensing

[GET MY OWN PROFILE](#)

TITLE	CITED BY	YEAR
Road segmentation of remotely-sensed images using deep convolutional neural networks with landscape metrics and conditional random fields T Panboonyuen, K Jitkajornwanich, S Lawawirojwong, P Srestasathien, ... Remote Sensing 9 (7), 680	137	2017
Semantic segmentation on remotely sensed images using an enhanced global convolutional network with channel attention and domain specific transfer learning T Panboonyuen, K Jitkajornwanich, S Lawawirojwong, P Srestasathien, ... Remote Sensing 11 (1), 83	107	2019
An enhanced deep convolutional encoder-decoder network for road segmentation on aerial imagery T Panboonyuen, P Vateekul, K Jitkajornwanich, S Lawawirojwong Recent Advances in Information and Communication Technology 2017 ...	47	2018
Transformer-based decoder designs for semantic segmentation on remotely sensed images T Panboonyuen, K Jitkajornwanich, S Lawawirojwong, P Srestasathien, ... Remote Sensing 13 (24), 5100	45	2021
Object detection of road assets using transformer-based YOLOX with feature pyramid decoder on thai highway panorama T Panboonyuen, S Thongbai, W Wongweeranimit, P Santitamnont, ... Information 13 (1), 5	20	2021
Real-time polyps segmentation for colonoscopy video frames using compressed fully convolutional network I Wichakarn, T Panboonyuen, C Udomcharoenchaikit, P Vateekul MultiMedia Modeling: 24th International Conference, MMM 2018, Bangkok ...	18	2018
Semantic segmentation on medium-resolution satellite images using deep convolutional networks with remote sensing derived indices S Chanharaj, K Pornratthanapong, P Chitsimpachayakun, T Panboonyuen, ... 2018 15th International joint conference on computer science and software ...	13	2018

Cited by

	All	Since 2019
Citations	429	411
h-index	9	9
i10-index	8	8



Public access [VIEW ALL](#)

0 articles [4 articles](#)

not available [available](#)

Based on funding mandates



Featured Publications

≡ Google Scholar

Q SIGN IN

Teerapong Panboonyuen

Senior Research Scientist at MARS, Post-doc at Chula
Verified email at chula.ac.th - [Homepage](#)

Artificial Intelligence Machine Learning Deep Learning Computer Vision Remote Sensing

[GET MY OWN PROFILE](#)

Cited by

	All	Since 2019
Citations	429	411
h-index	9	9
i10-index	8	8

TITLE	CITED BY	YEAR
Road segmentation of remotely-sensed images using deep convolutional neural networks with landscape metrics and conditional random fields	137	2017

T Panboonyuen, K Jitkajornwichech, S Lawawirojwong, P Srestasathien, ...
Remote Sensing 9 (7), 680

Semantic segmentation on remotely sensed images using an enhanced global convolutional network with channel attention and domain specific transfer learning	107	2019
---	-----	------

T Panboonyuen, K Jitkajornwichech, S Lawawirojwong, P Srestasathien, ...
Remote Sensing 11 (1), 83

An enhanced deep convolutional encoder-decoder network for road segmentation on aerial imagery	47	2018
--	----	------

T Panboonyuen, P Vateekul, K Jitkajornwichech, S Lawawirojwong
Recent Advances in Information and Communication Technology 2017, ...

Transformer-based decoder designs for semantic segmentation on remotely sensed images	45	2021
---	----	------

T Panboonyuen, K Jitkajornwichech, S Lawawirojwong, P Srestasathien, ...
Remote Sensing 13 (24), S100

Object detection of road assets using transformer-based YOLOX with feature pyramid decoder on Thai highway panorama	20	2021
---	----	------

P Panichkitkosolkul, NV Wongweewarnit, P Sanitannont, ...
Information 13 (1), 5

Real-time polyps segmentation for colonoscopy video frames using compressed fully convolutional network	18	2018
---	----	------

N Ichikami, T Panboonyuen, C Udomcharoenchaikit, P Vateekul
MultiMedia Modeling 24th International Conference, MWM 2018, Bangkok, ...

Semantic segmentation on medium-resolution satellite images using deep convolutional networks with remote sensing derived indices	13	2018
---	----	------

S Chanthana, A Ponnattanapong, P Chinsupayakun, T Panboonyuen, ...
2018 15th International joint conference on computer science and software ...

Public access

0 articles	4 articles
------------	------------

not available

available

Based on funding mandates

What Is Computer Vision? [Basic Tasks & Techniques]


Semantic
Segmentation



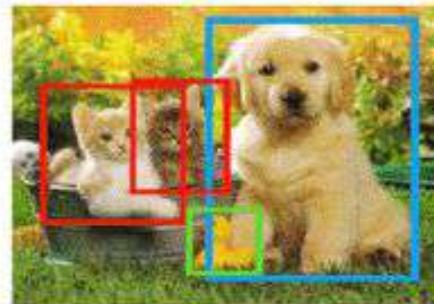
Classification



Classification
+ localization



Object detection



Instance
segmentation



CAT GRASS
TREE

CAT

CAT

CAT DOG DUCK

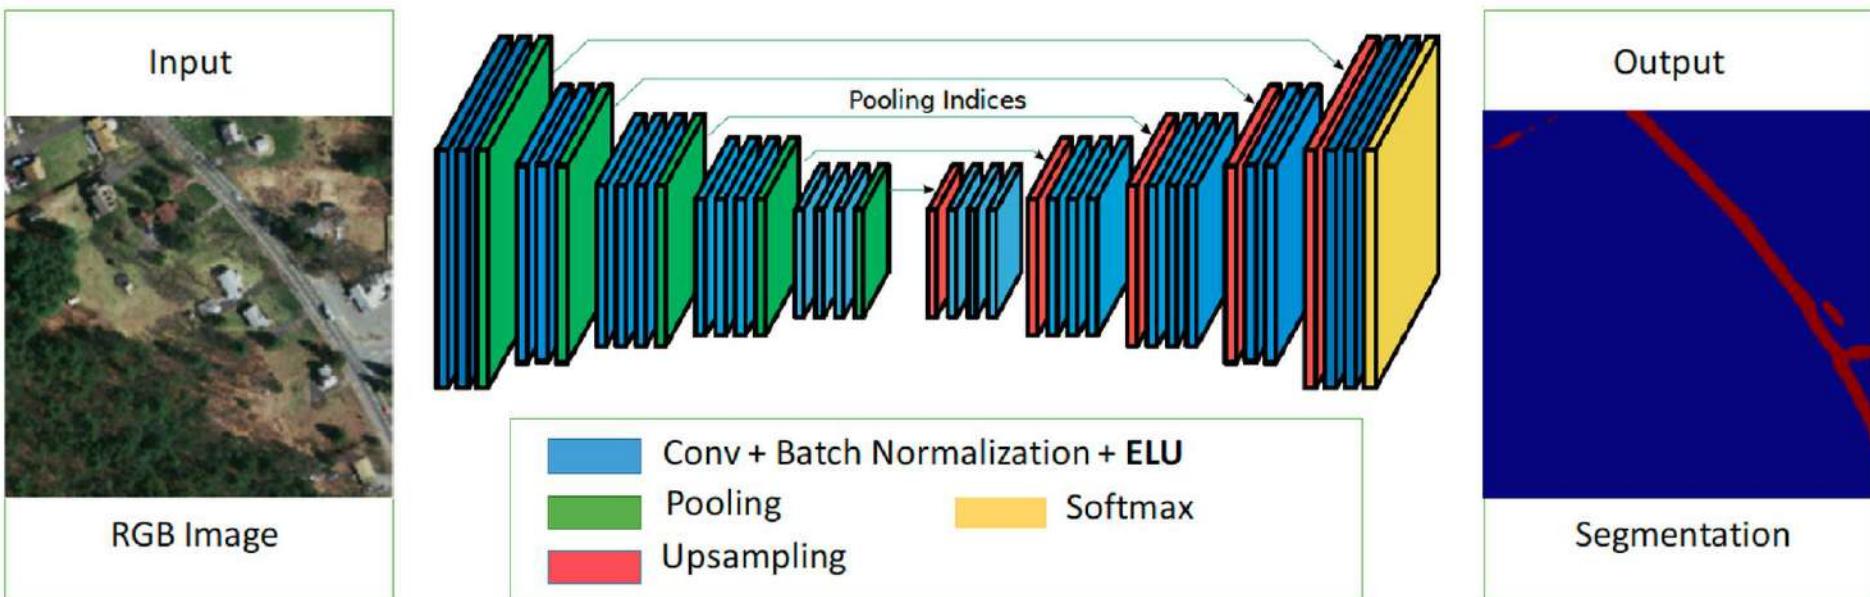
CAT CAT DOG DUCK

No object
Just pixels

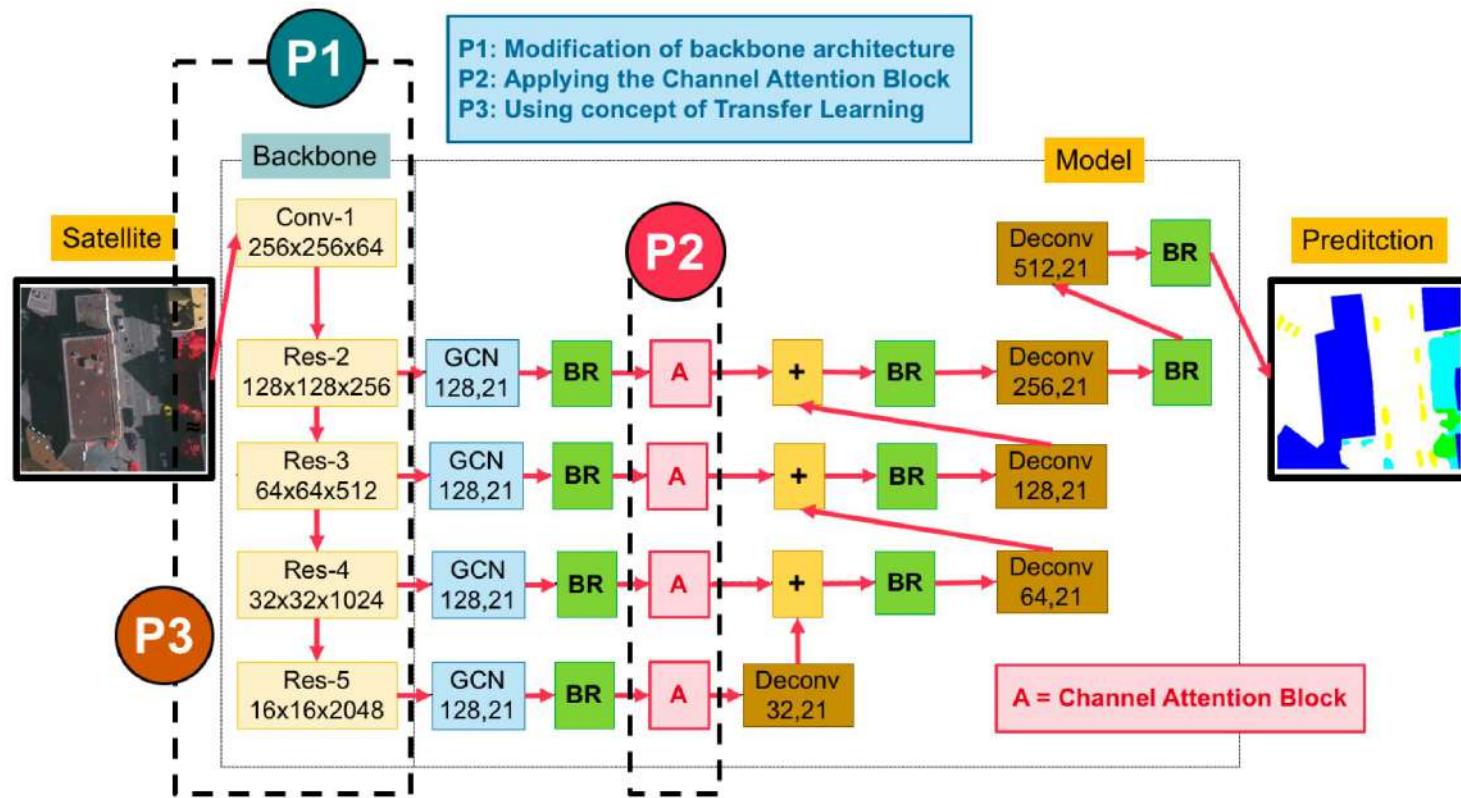
Single object

Multiple objects


[1] Panboonyuen, Teerapong, et al. "Road segmentation of remotely-sensed images using deep convolutional neural networks with landscape metrics and conditional random fields." *Remote Sensing* 9.7 (2017): 680.



[2] Panboonyuen, Teerapong, et al. "Semantic segmentation on remotely sensed images using an enhanced global convolutional network with channel attention and domain specific transfer learning." Remote Sensing 11.1 (2019): 83.

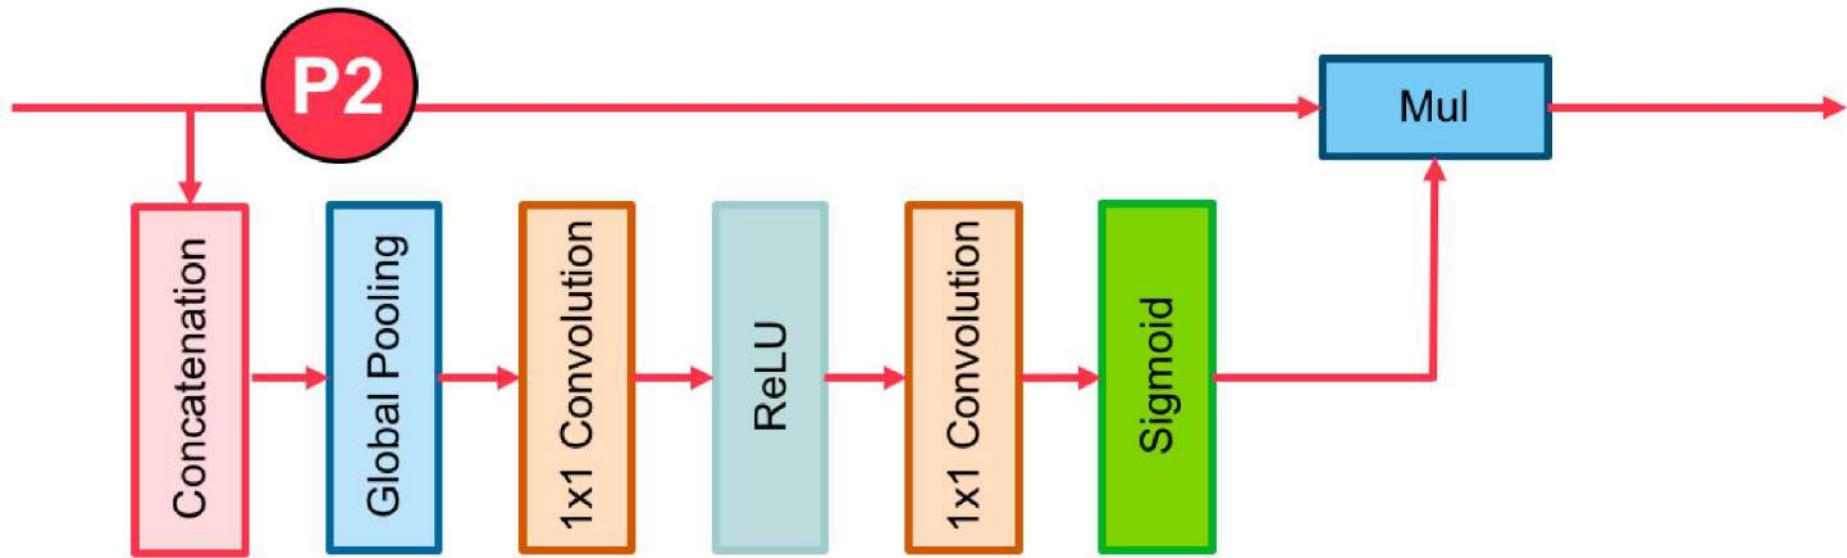


[2] Panboonyuen, Teerapong, et al. "Semantic segmentation on remotely sensed images using an enhanced global convolutional network with channel attention and domain specific transfer learning." Remote Sensing 11.1 (2019): 83.



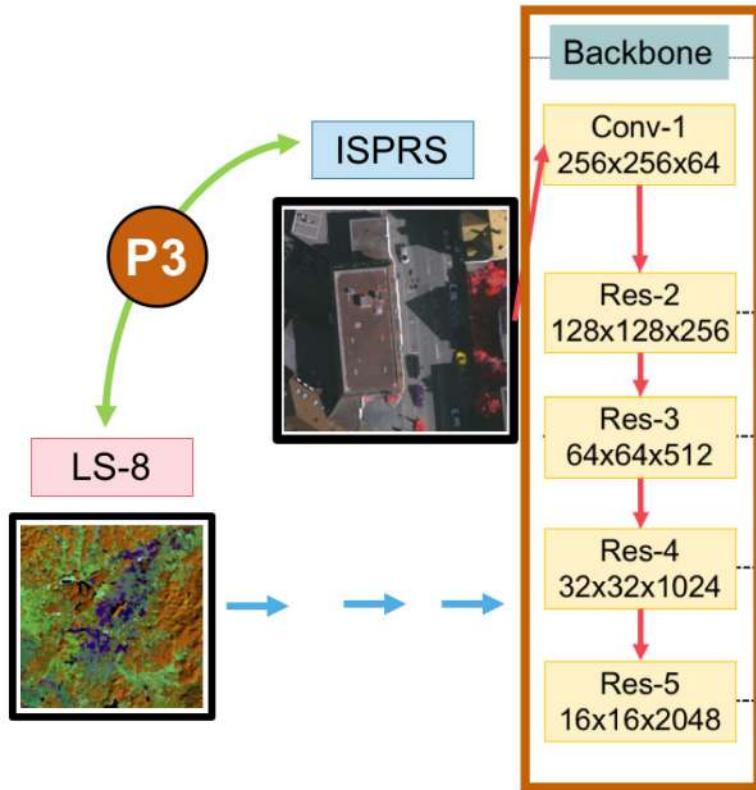
An overview of the whole backbone pipeline in (left) the main backbone with varying by ResNet50, ResNet 101, and ResNet 152; (right) the major drivers of my main classification network (composed of a global convolutional network (GCN) and a boundary refinement (BR) block).

[2] Panboonyuen, Teerapong, et al. "Semantic segmentation on remotely sensed images using an enhanced global convolutional network with channel attention and domain specific transfer learning." Remote Sensing 11.1 (2019): 83.



- Components of the channel attention block.
- The red lines represent the downsample operators, respectively.
- The red line cannot change the size of feature maps.
- It is only a path for information passing.

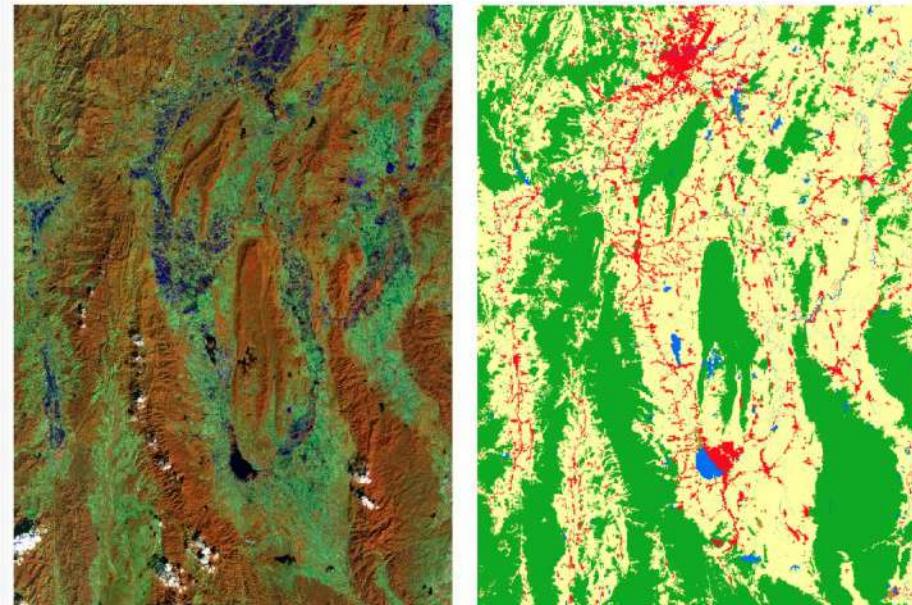
[2] Panboonyuen, Teerapong, et al. "Semantic segmentation on remotely sensed images using an enhanced global convolutional network with channel attention and domain specific transfer learning." Remote Sensing 11.1 (2019): 83.



- The domain-specific transfer learning strategy reuses pre-trained weights of models between two datasets—very high (ISPRS) and medium (Landsat-8; LS-8) resolution images.

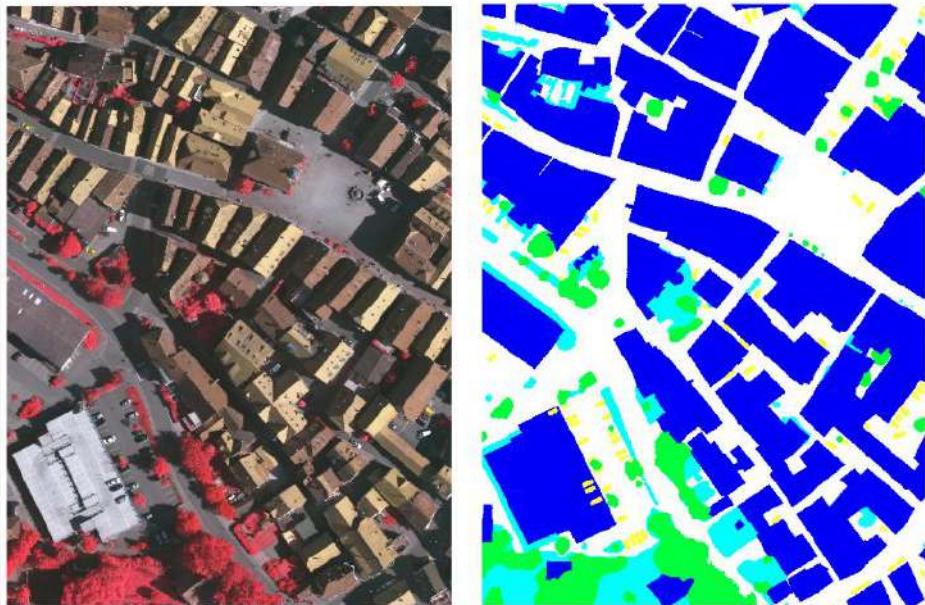
[2] Panboonyuen, Teerapong, et al. "Semantic segmentation on remotely sensed images using an enhanced global convolutional network with channel attention and domain specific transfer learning." Remote Sensing 11.1 (2019): 83.

Figure 6. Sample satellite images from Nan, a province in Thailand (**left**), and corresponding ground truth (**right**). The label of medium resolution dataset includes five categories: agriculture (yellow), forest (green), miscellaneous (brown), urban (red), and water (blue).



X

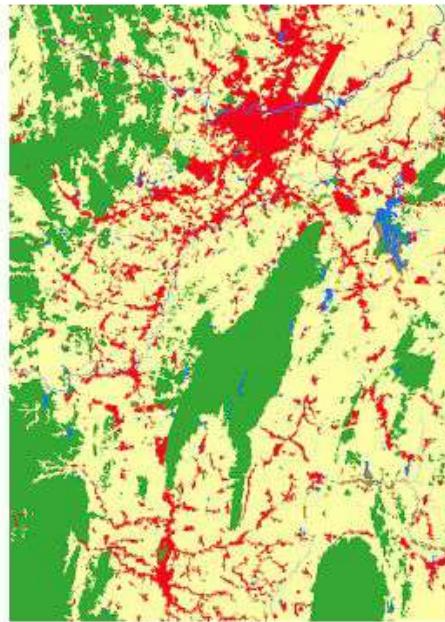
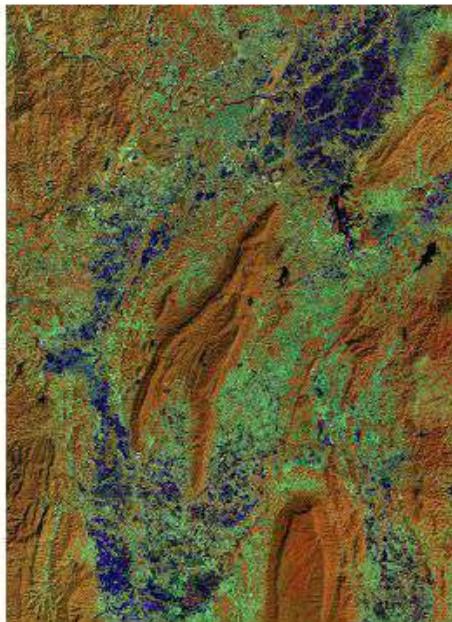
Figure 8. The sample input tile from Figure 7 (**left**) and corresponding ground truth (**right**). The label of the Vaihingen Challenge includes six categories: impervious surface (imp surf, white), building (blue), low vegetation (low veg, cyan), tree (green), car (yellow), and clutter/background (red).



X

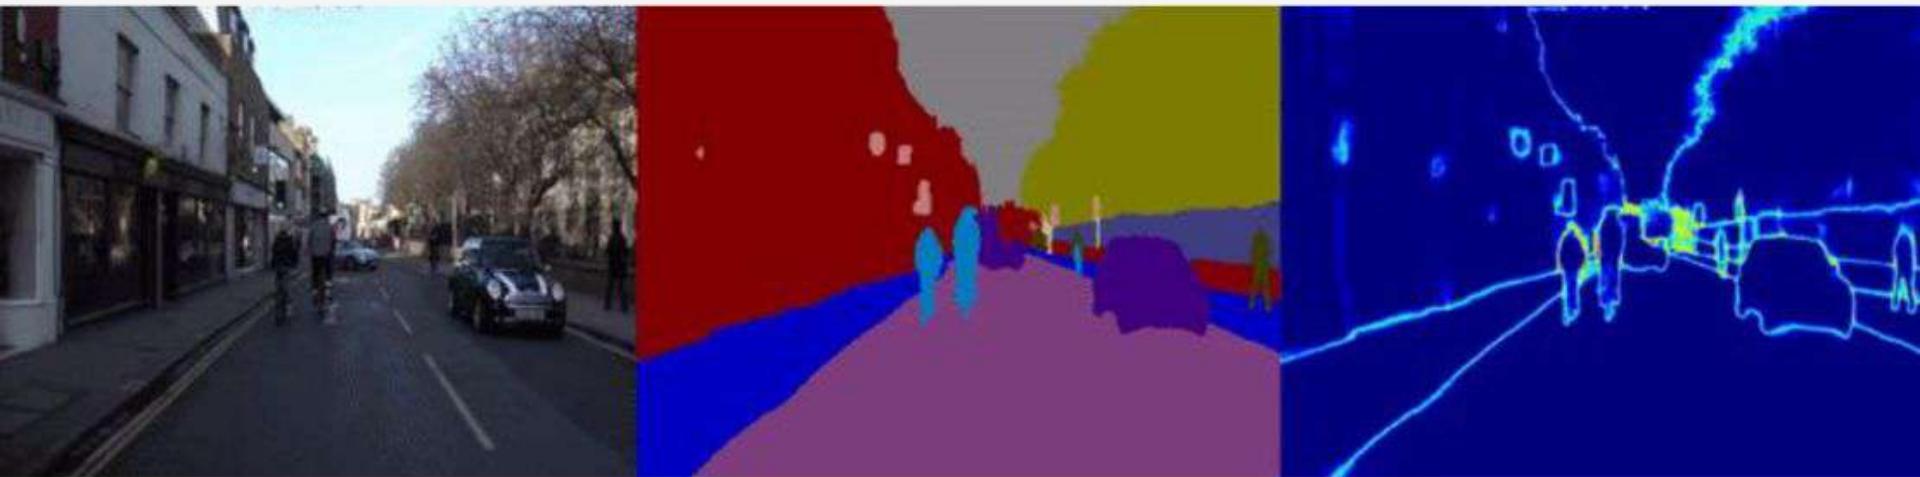
Public and Private Corpora

Private corpus (GISTDA Nan Province Corpus)



Color	Class
Yellow	Agriculture
Green	Forest
Brown	Miscellaneous
Red	Urban
Blue	Water

Bayesian SegNet for probabilistic scene understanding




Input Image

Semantic Segmentation

Uncertainty

What kind of uncertainty can we model?

Epistemic uncertainty is *modeling uncertainty*
Aleatoric uncertainty is *sensing uncertainty*



Modeling Uncertainty with Bayesian Deep Learning



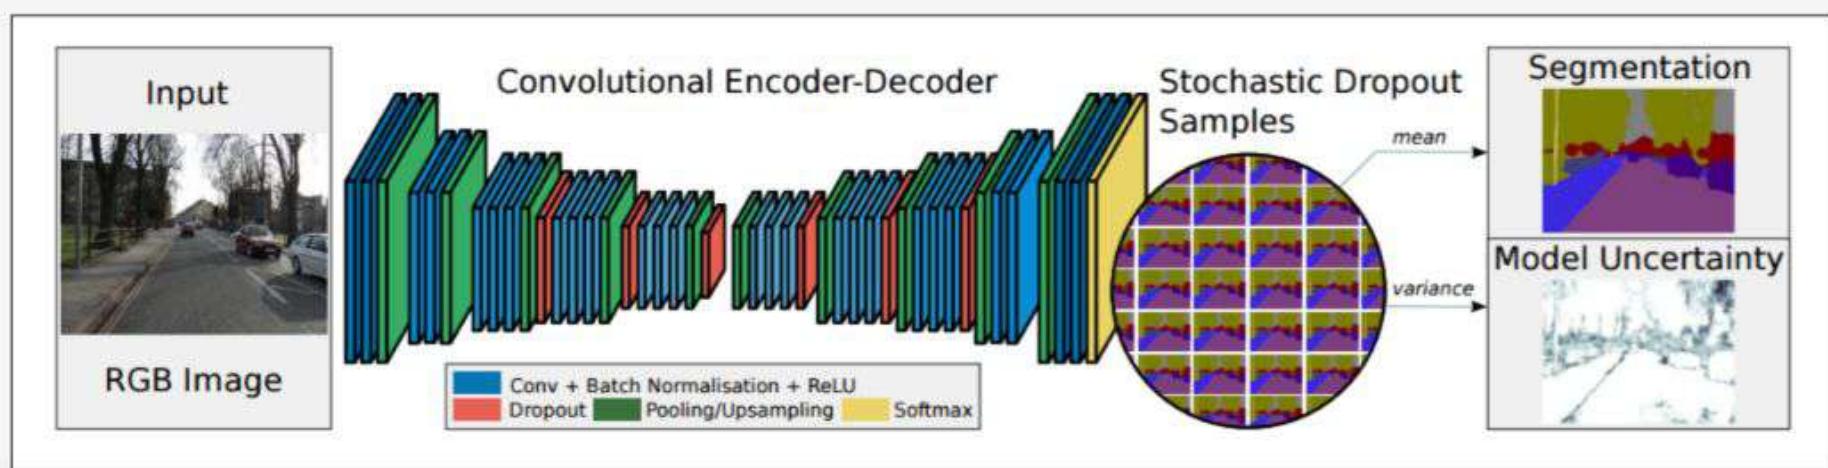
Deep learning is required to achieve state of the art results in computer vision applications but doesn't provide uncertainty estimates.

- **Bayesian neural networks** are a framework for understanding uncertainty in deep learning
- They have **distributions over network parameters** (rather than deterministic weights)
- Traditionally they have been **tricky to scale**

Modeling Epistemic Uncertainty with Bayesian Deep Learning

We can **model epistemic uncertainty** in deep learning models using Monte Carlo **dropout sampling** at test time.

Dropout sampling can be interpreted as **sampling from a distribution over models**.



Modeling Aleatoric Uncertainty with Probabilistic Deep Learning

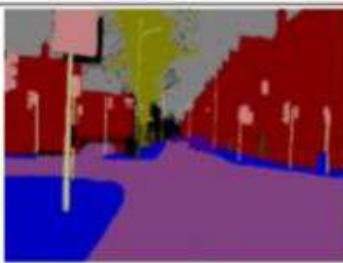
	Deep Learning	Probabilistic Deep Learning
Model	$[\hat{y}] = f(x)$	$[\hat{y}, \hat{\sigma}^2] = f(x)$
Regression	$Loss = \ y - \hat{y}\ ^2$	$Loss = \frac{\ y - \hat{y}\ ^2}{2\hat{\sigma}^2} + \log \hat{\sigma}^2$
Classification	$Loss = SoftmaxCrossEntropy(\hat{y}_t)$	$\hat{y}_t = \hat{y} + \epsilon_t \quad \epsilon_t \sim N(0, \hat{\sigma}^2)$ $Loss = \frac{1}{T} \sum_t SoftmaxCrossEntropy(\hat{y}_t)$

Semantic Segmentation Performance on CamVid

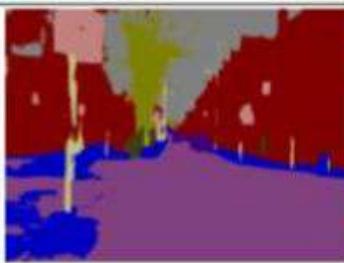
CamVid Results	IoU Accuracy
DenseNet (State of the art baseline)	67.1
+ Aleatoric Uncertainty	67.4
+ Epistemic Uncertainty	67.2
+ Aleatoric & Epistemic	67.5



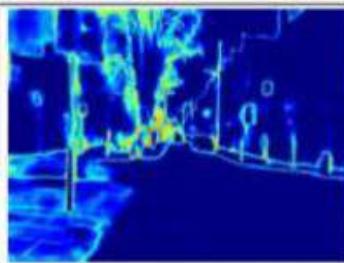
(a) Input Image



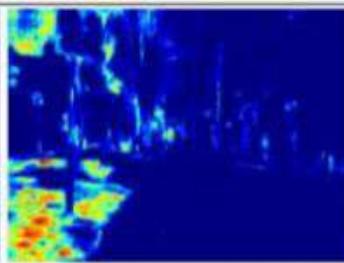
(b) Ground Truth



(c) Semantic Segmentation

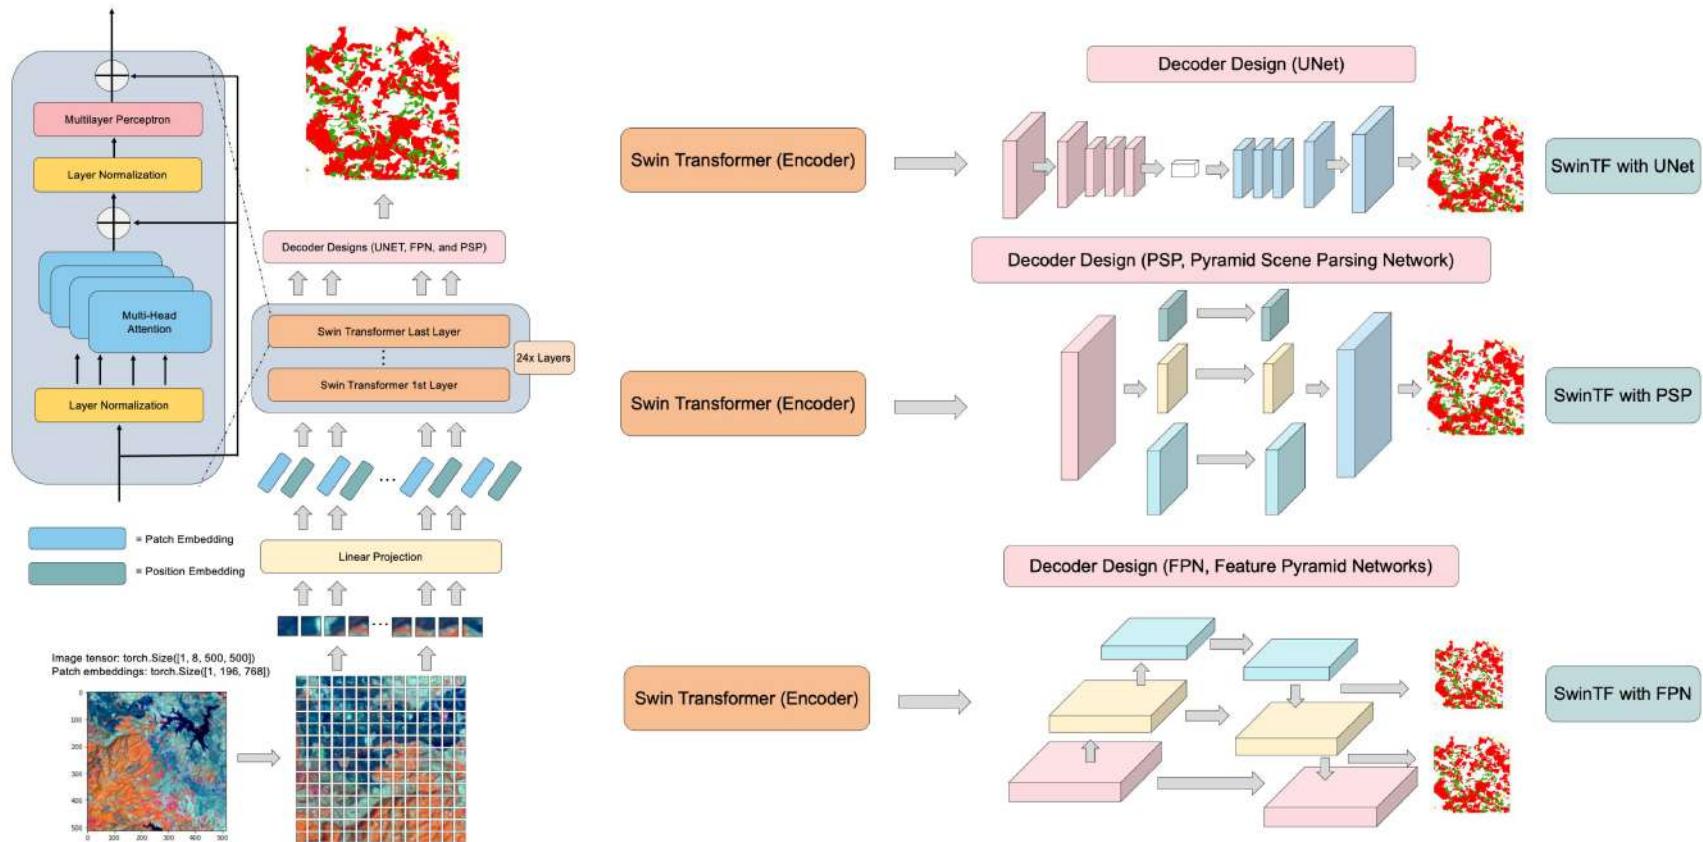


(d) Aleatoric Uncertainty



(e) Epistemic Uncertainty

[3] Panboonyuen, Teerapong, et al. "Transformer-based decoder designs for semantic segmentation on remotely sensed images." Remote Sensing 13.24 (2021): 5100.



[4] Wichakam, I., Panboonyuen, T., Udomcharoenchaikit, C., & Vateekul, P. (2018). **Real-time polyps segmentation for colonoscopy video frames using compressed fully convolutional network**. In MultiMedia Modeling: 24th International Conference, MMM 2018, Bangkok, Thailand, February 5-7, 2018, Proceedings, Part I 24 (pp. 393-404). Springer International Publishing.

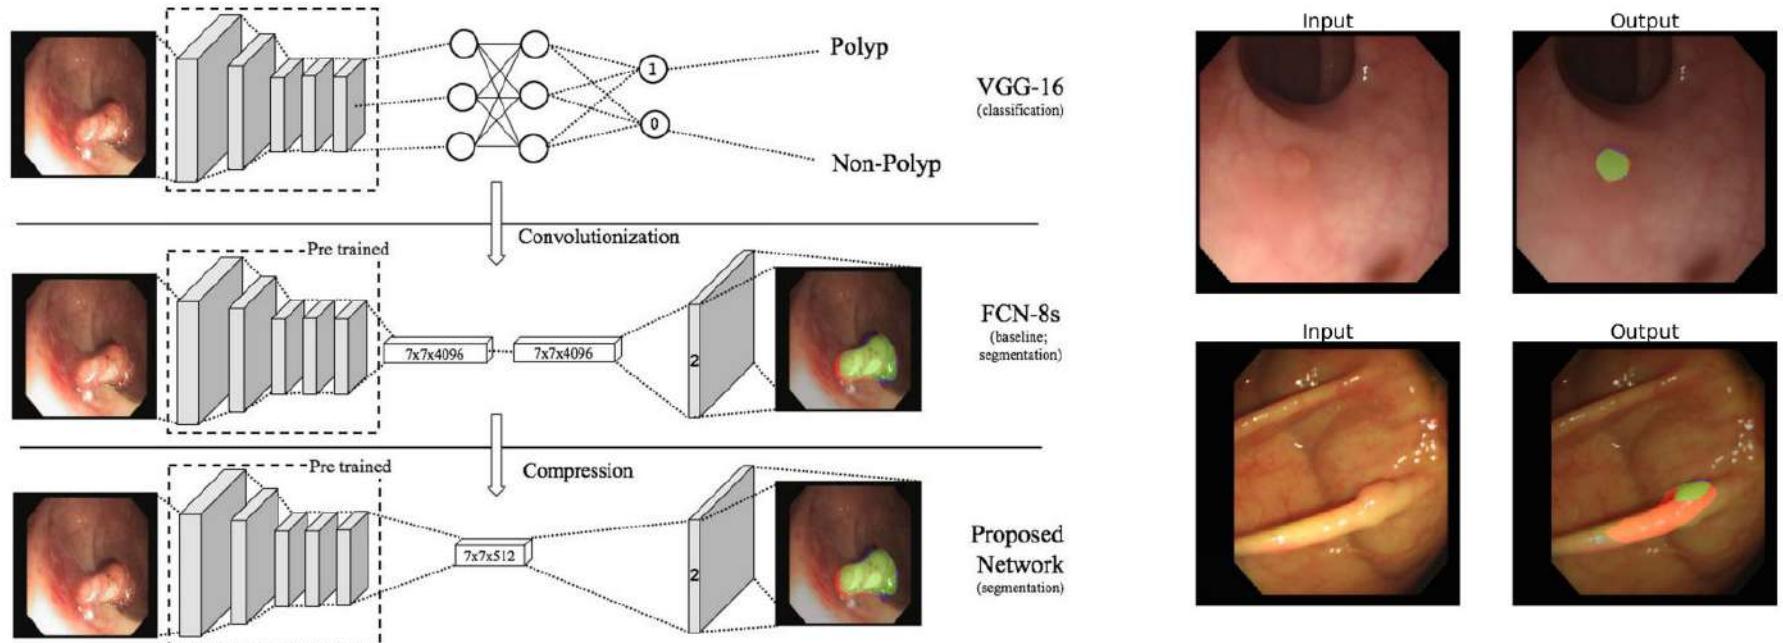




Fig. 1. Overview of our compressed network which is compressed from the original FCN-8s [8] based on VGG-16 [4] architecture.

[5] Panboonyuen, Teerapong, et al. "Object detection of road assets using transformer-based YOLOX with feature pyramid decoder on thai highway panorama." Information 13.1 (2022): 5.



[6] Thitisiriwech, K., Panboonyuen, T., Kantavat, P., Iwahori, Y., & Kijsirikul, B. (2022). The Bangkok Urbanscapes Dataset for Semantic Urban Scene Understanding Using Enhanced Encoder-Decoder With Atrous Depthwise Separable A1 Convolutional Neural Networks. *IEEE Access*, 10, 59327-59349.




FIGURE 6. An overview of enhanced DeepLab-V3+ (Encoder-Decoder with atrous separable convolutional for semantic segmentation [15]) with ResNet-101 backbone [41] (DeepLab-V3-A1-ResNet-101).




FIGURE 8. Sample 1: The example of Sukhumvit's large road from the training set of the Bangkok Urbanscapes dataset. The input image is shown in (a), and the ground truth is shown in (b).

Void	Building	Wall	Tree	VegetationMisc	Fence
Sidewalk	ParkingBlock	Column_Pole	TrafficCone	Bridge	SignSymbol
Misc_Text	TrafficLight	Sky	Tunnel	Archway	Road
ReadShoulder	LaneMrgsDriv	LaneMrgsNonDriv	Animal	Pedestrian	Child
CarLuggagePram	Bicyclist	MotorcycleBcooter	Car	SUV/PickupTruck	Truck_Bus
Train	OtherMoving				

FIGURE 3. The semantic color codes of the CamVid dataset. Each color is encoded with respect to the semantic class in the ground truth images.

[6] Vajeethaveesin, T., Panboonyuen, T., Lawawironjwong, S., Srestasathiern, P., Jaiyen, S., & Jitkajornwanich, K. (2022). **A performance comparison between GIS-based and neuron network methods for flood susceptibility assessment in ayutthaya province.** Trends in Sciences, 19(2), 2038-2038.




Figure 2 The historical flood zone shows flood events, which were captured in 2016 from GISTDA. Most of the flooding areas are located in the west of the study area.




Figure 4 Flood susceptibility map from the flood risk assessment model (FRAM).

Classification of Primary Angle Closure Glaucoma (PACG) Using Deep Convolutional Neural Networks




Table 4: Performance comparison on F1 score with existing models on the WTW-LAF23 dataset.

Model	F1 score			
	Mean	SD	Min	Max
VGG19 [Dutta et al., 2016]	61.34	4.12	54.67	65.76
DenseNet121 [Targ et al., 2016]	61.67	4.01	54.89	65.91
RestNet50 [Huang et al., 2017]	65.93	6.11	56.11	76.33
Transformer [Liu et al., 2021]	58.18	5.07	51.34	69.83
Our Proposed Method (SimCLR/Chen et al. [2020])	67.82	N/A	N/A	N/A



Table 1: Details of used visual datasets.

Dataset	Type	Positive Size	Negative Size	Train Size	Val Size	# Classes
LAFx10_spss_Frequency2.3	Classification	612	612	1,101	123	2
AQD_Frequency2.4	Classification	3,338	3,153	5,841	650	2

Table 2: Details of used visual datasets.

Dataset	Type	Female Size	Male Size	Unknown	Other	Train Size	Val Size	# Classes
GENDER	Classification	3,257	2,106	25	1	3,754	804	2

Table 3: Details of used visual datasets.

Dataset	Type	Young Adults	Middle-Aged	Old Adults	Train Size	Val Size	# Classes
AGE	Classification	224	3,885	1,271	3,765	807	3

มาร์สคือผู้นำในการพัฒนาระบบนิเวศสำหรับผู้ใช้รถยนต์ ปลอดภัยจากการตรวจสภาพรถยนต์แบบเดิมด้วยการใช้ AI ให้การตรวจสภาพรถยนต์เป็นเรื่องที่ง่ายกว่า

มาร์สสามารถเชื่อมโยงธุรกิจในหลากหลายอุตสาหกรรมกับผู้ใช้รถยนต์ให้เกิดความสะดวกและรวดเร็วยิ่งขึ้น มาร์สหรือ MARS (Motor AI Recognition Solution) เป็นแอปพลิเคชันที่ปลอดภัยในการตรวจสภาพรถยนต์แบบเดิมๆ ด้วยเทคโนโลยี AI แบบครบวงจรที่มาร์สได้คิดค้นและพัฒนาขึ้นมาเพื่อช่วยให้ธุรกิจของคุณลดต้นทุนด้านบุคลากร ลดขั้นตอนการทำงาน เพิ่มความรวดเร็วและแม่นยำในการตรวจสอบและจัดเก็บข้อมูล

[ดูโซลูชันทั้งหมด](#)[เกี่ยวกับมาร์ส](#)

<https://www.marssolution.io/>

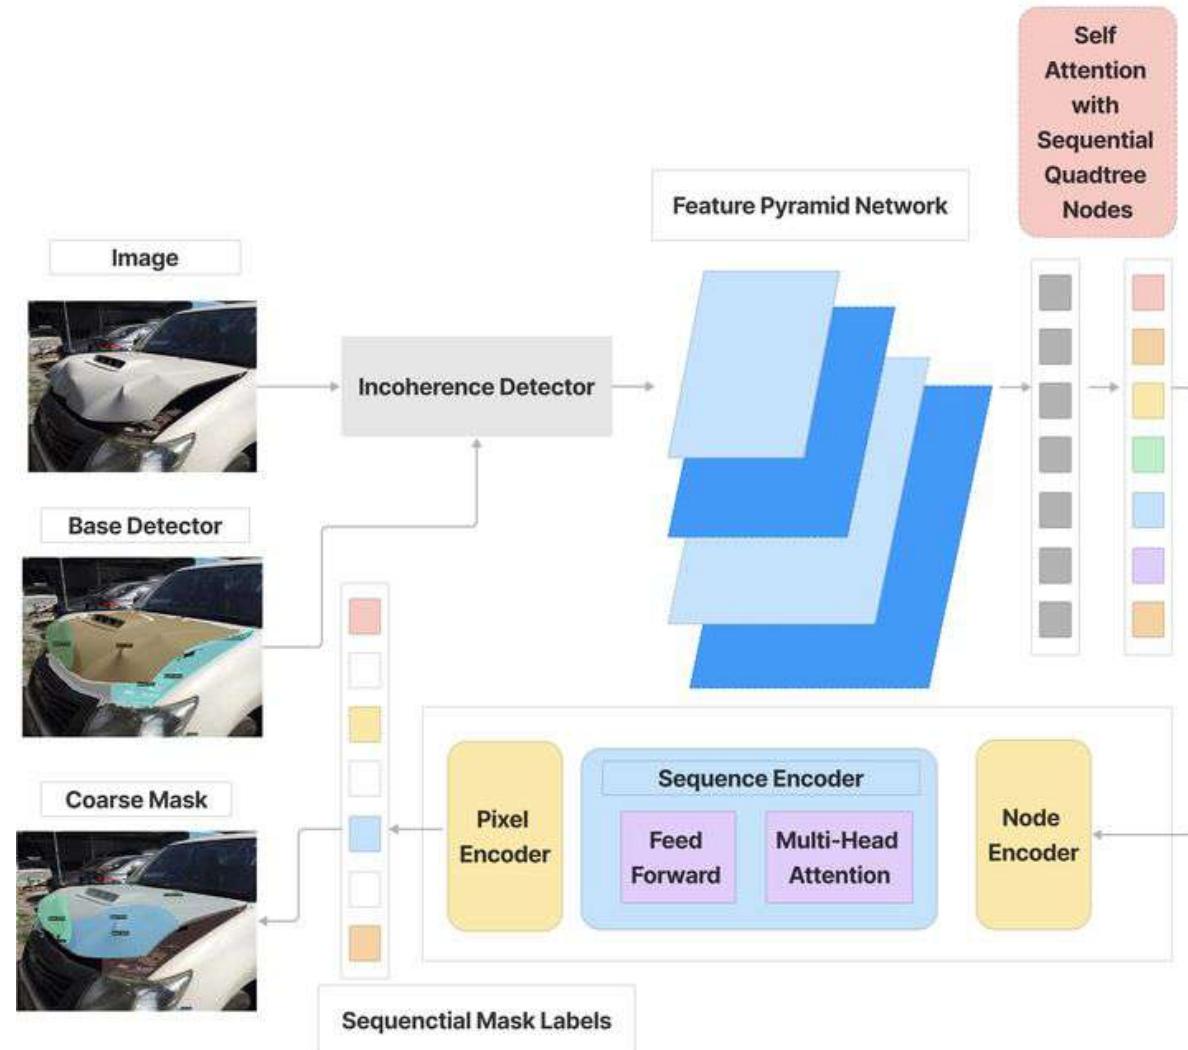
ดาวน์โหลด MARS Inspect หรือ “มาตรวจ” วันนี้


ให้เรื่องการตรวจรถยนต์เป็นเรื่องง่ายในมือคุณ เพียงแค่ โหลด เตรียม ถ่าย





Panboonyuen,
Teerapong, et al.
**"MARS: Mask
Attention Refinement
with Sequential
Quadtree Nodes for
Car Damage Instance
Segmentation."**
*International
Conference on Image
Analysis and
Processing.* Cham:
Springer Nature
Switzerland, 2023.





Large Language Model (LLM)

[*'lärj 'laŋ-gwij 'mä-dəl*]

A deep learning algorithm that's equipped to summarize, translate, predict, and generate human-sounding text to convey ideas and concepts.



HAVE YOU FIGURED
OUT HOW AI WILL
IMPACT OUR
BUSINESS?

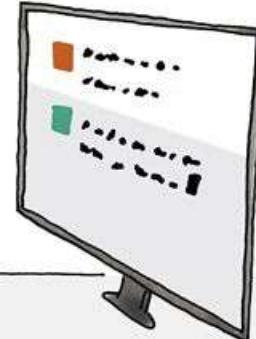
WORKING
ON IT.



How will AI impact
our business?



There are many ways
that AI can impact ■



TOM
FISH
BURNE

Understanding Multimodal



Chip Huyen

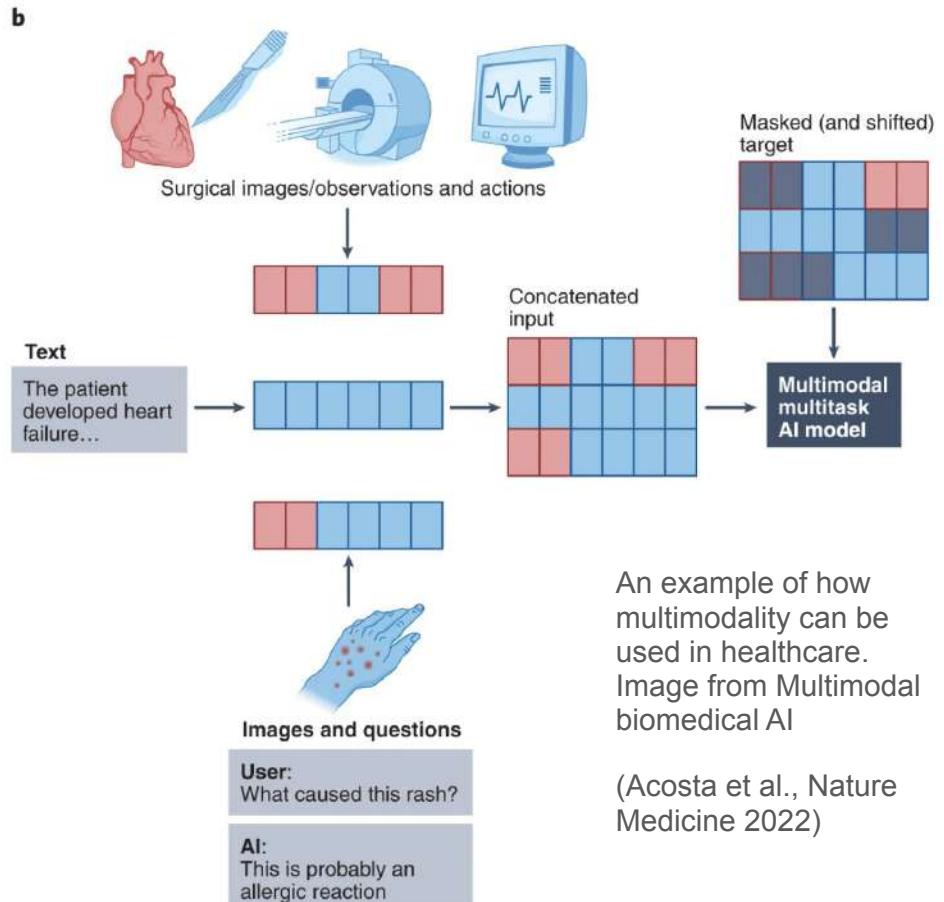
@chipro

New blog post: Multimodality and Large Multimodal Models (LMMs)

Being able to work with data of different modalities -- e.g. text, images, videos, audio, etc. -- is essential for AI to operate in the real world.

This post covers multimodal systems in general, including Large Multimodal Models. It consists of 3 parts.

- * Part 1 covers the context for multimodality.
- * Part 2 discusses how to train a multimodal system, using the architectures of CLIP and Flamingo, and examples from GPT-4V.
- * Part 3 discusses some active research areas for LMMs, including generating multimodal outputs.




An example of how multimodality can be used in healthcare.
Image from Multimodal biomedical AI

(Acosta et al., Nature Medicine 2022)

<https://huyenchip.com/2023/10/10/multimodal.htm>

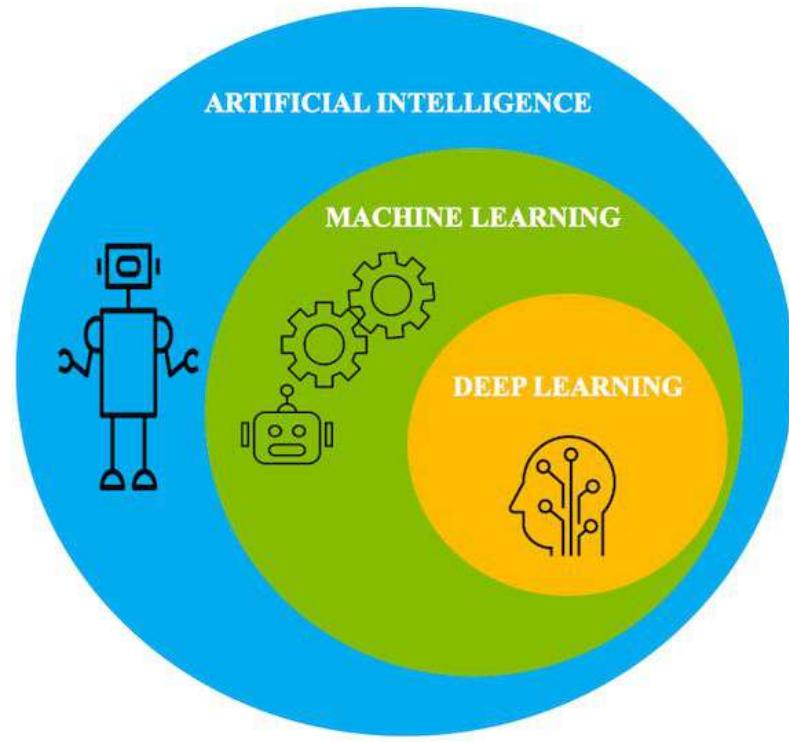
<https://twitter.com/chipro/status/1711970025874321479/photo/1>

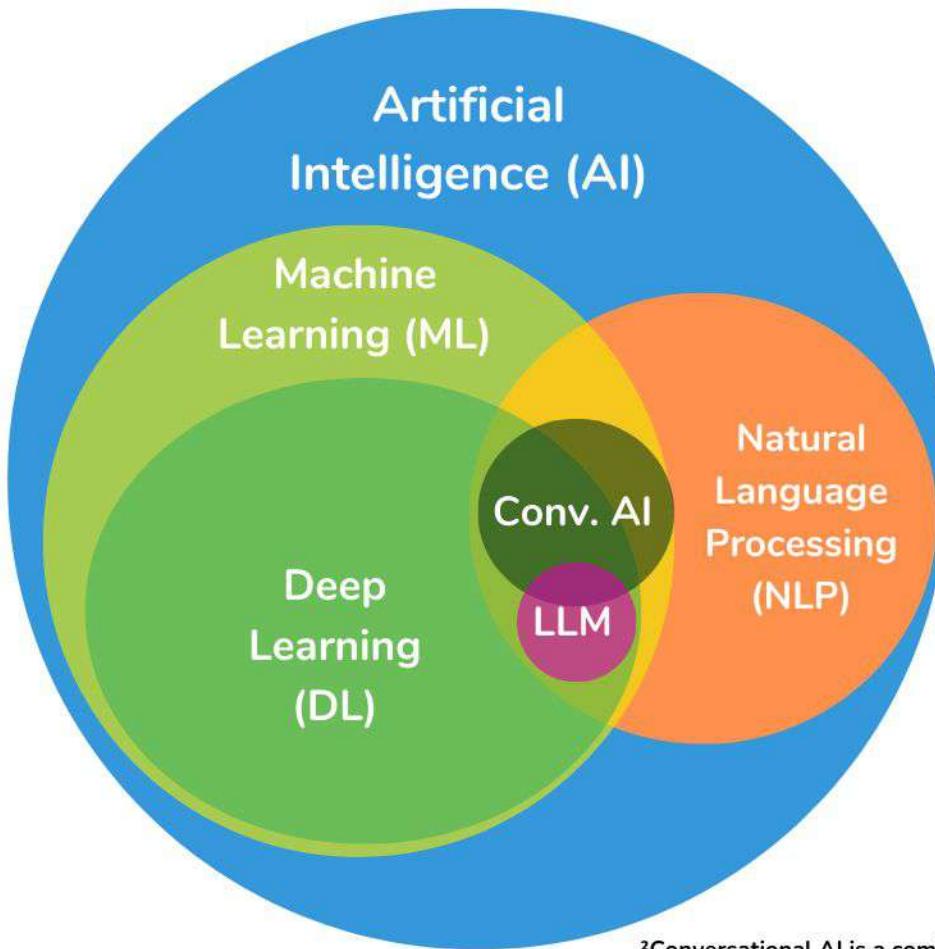
Basic AI



AI vs Machine Learning vs Deep Learning

- **Artificial Intelligence (AI)** is the concept of creating smart intelligent machines.
- **Machine Learning (ML)** is a subset of artificial intelligence that helps you build AI-driven applications.
- **Deep Learning** is a subset of machine learning that uses vast volumes of data and complex algorithms to train a model.






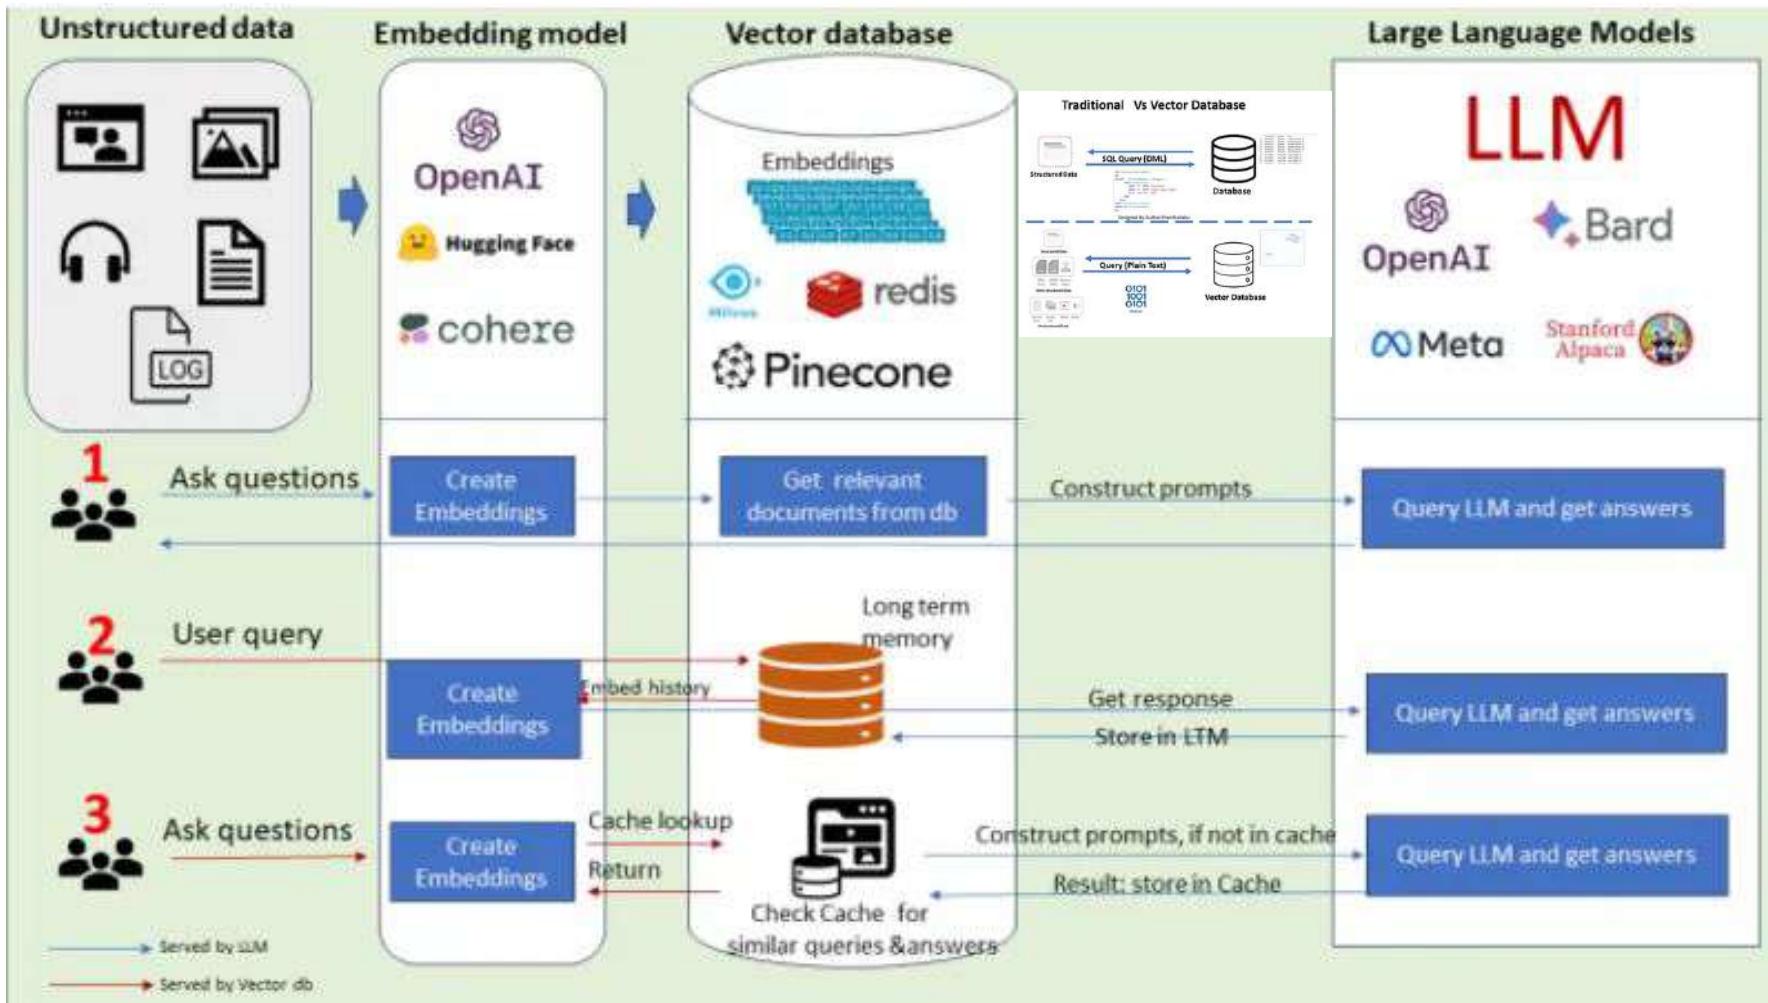
- Artificial Intelligence (AI)
- Machine Learning (ML)
- Deep Learning (DL)
- Natural Language Processing (NLP)
- Large Language Model (LLM)¹
- Conversational AI (Conv. AI)²

¹LLM is an intersection of DL and NLP

²Conversational AI is a combination of ML and NLP. It may include DL and LLM, but that isn't always the case.

What Is Generative AI?





Difference between Artificial Intelligence and Artificial General Intelligence (AGI)

Artificial Intelligence (AI)

- AI mimics human intelligence.
- AI is trained by data scientists on specific, singular, or a limited number of tasks.
Examples: image recognition, speech-to-text transcription, text generation, chatbots.
- Is not self-aware; has no consciousness or ability to think on its own.

Artificial General Intelligence (AGI)

- AGI equals human intelligence, theoretically.
- AGI possesses common sense and creativity and expresses emotions.
- Has the ability to learn, apply knowledge, generalize, and plan ahead.
- Not fully realized yet. Some doubt if it ever will be.

Machine Learning



Narrow Artificial Intelligence (ANI)

Stage One: Machines imitate human behavior, specializing in one area to solve a problem.

i.e. Siri, ChatGPT, Alexa


Machine Intelligence



Artificial General Intelligence (AGI)

Stage Two: Machines can continuously learn and are as smart as humans.

Machine Consciousness



Artificial Super Intelligence (ASI)

Stage Three: Machines that are smarter than humans across the board.

Meta scientist Yann LeCun says AI won't destroy jobs forever

15 June 2023

Share 

Chris Vallance
Technology reporter



Prof Yann LeCun is known as one of the three godfathers of AI and works as Facebook-owner Meta's top AI scientist



Yann LeCun Criticises Elon Musk for ‘Batsh*t-Crazy’ Conspiracy Theories

ENDLESS ORIGINS

PIM

Forward-Forward Algorithm: Will it replace Backpropagation?

The Forward-Forward Algorithm: Some Preliminary Investigations

Geoffrey Hinton
Google Brain
geoffhinton@google.com

Abstract

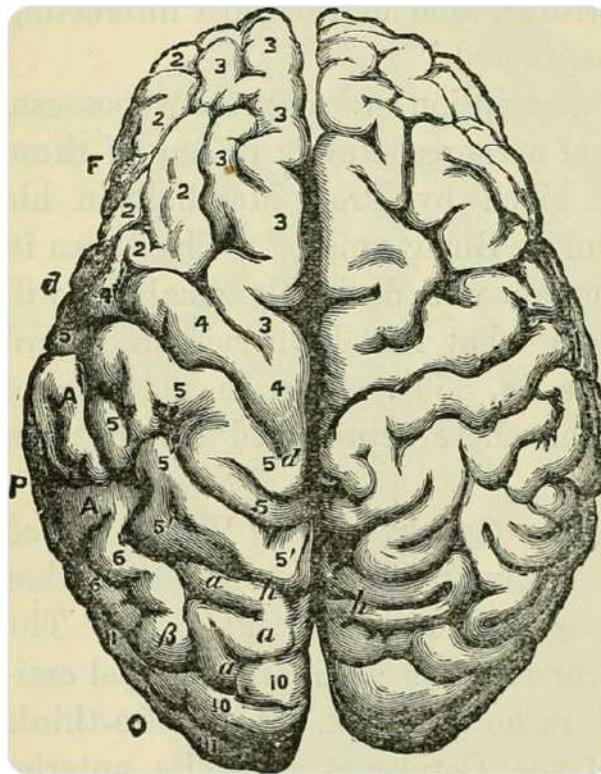
The aim of this paper is to introduce a new learning procedure for neural networks and to demonstrate that it works well enough on a few small problems to be worth serious investigation. The Forward-Forward algorithm replaces the forward and backward passes of backpropagation by two forward passes, one with positive (*i.e.* real) data and the other with negative data which could be generated by the network itself. Each layer has its own objective function which is simply to have high goodness for positive data and low goodness for negative data. The sum of the squared activities in a layer can be used as the goodness but there are many other possibilities, including minus the sum of the squared activities. If the positive and negative passes can be separated in time, the negative passes can be done offline, which makes the learning much simpler in the positive pass and allows video to be pipelined through the network without ever storing activities or stopping to propagate derivatives.





Martin Görner @martin_gorner · Dec 5, 2022

I seems very unlikely that the human brain uses back propagation to learn. There is little evidence of backprop mechanics in biological brains (no error derivatives propagating backwards, no storage of neuron activities to use in a packprop pass, ...).



...



Martin Görner

@martin_gorner

Also, the brain can learn from a continuous stream of incoming data and does not need to stop to run a backprop pass. Yes, sleep is beneficial for learning somehow, but we can learn awake too.



...

The Illustrated Transformer

Attention Is All You Need

Ashish Vaswani*
Google Brain
avaswani@google.com

Noam Shazeer*
Google Brain
noam@google.com

Niki Parmar*
Google Research
nikip@google.com

Jakob Uszkoreit*
Google Research
usz@google.com

Llion Jones*
Google Research
llion@google.com

Aidan N. Gomez* †
University of Toronto
aidan@cs.toronto.edu

Lukasz Kaiser*
Google Brain
lukaszkaiser@google.com

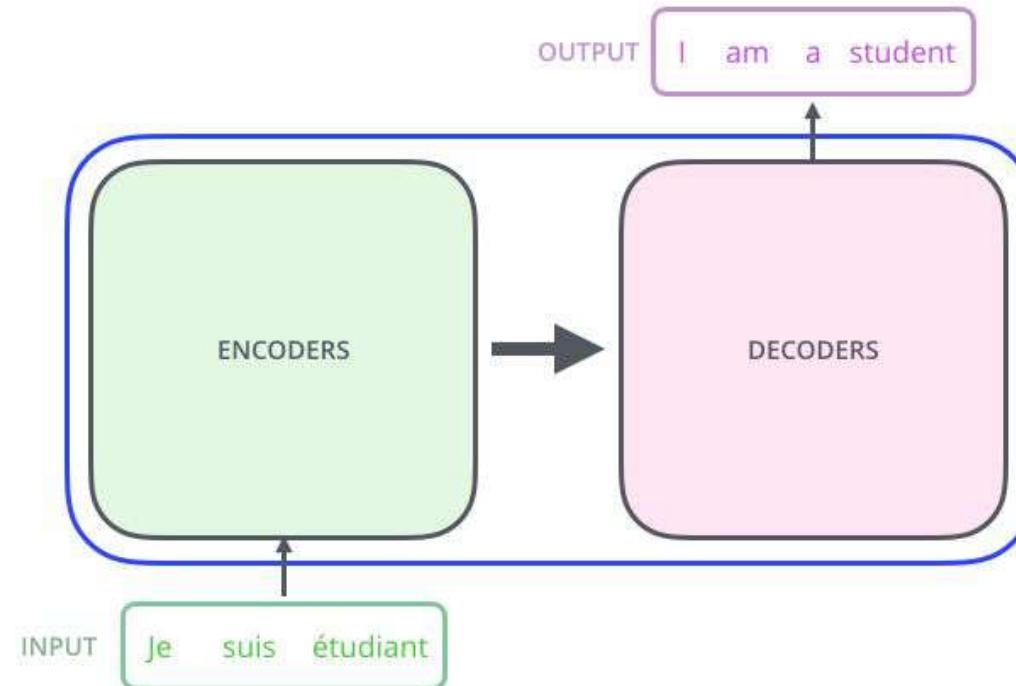
Illia Polosukhin* ‡
illia.polosukhin@gmail.com



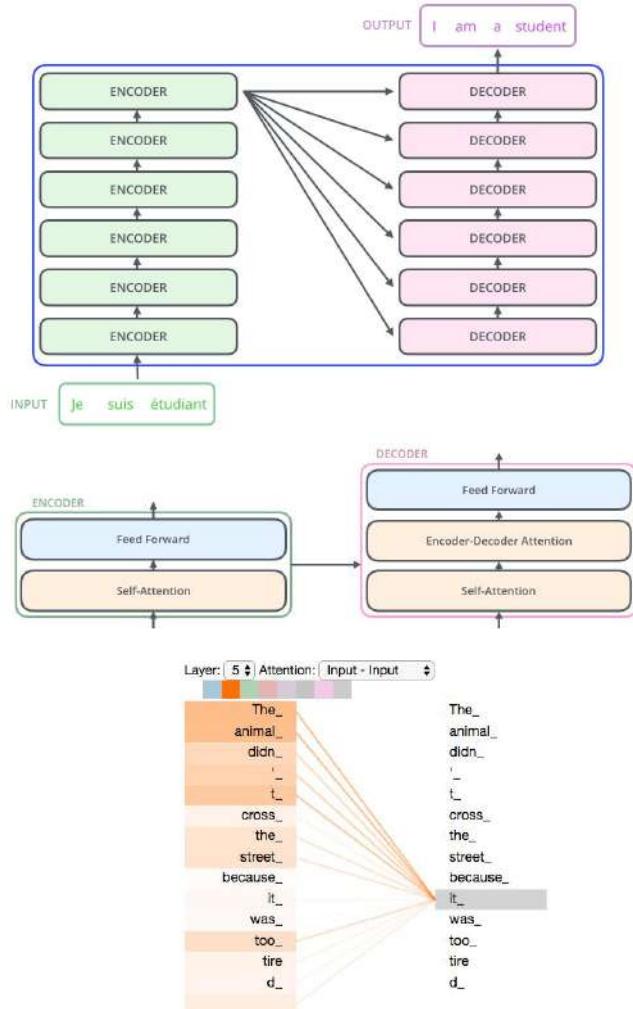

Abstract

The dominant sequence transduction models are based on complex recurrent or convolutional neural networks that include an encoder and a decoder. The best performing models also connect the encoder and decoder through an attention mechanism. We propose a new simple network architecture, the Transformer, based solely on attention mechanisms, dispensing with recurrence and convolutions entirely. Experiments on two machine translation tasks show these models to be superior in quality while being more parallelizable and requiring significantly less time to train. Our model achieves 28.4 BLEU on the WMT 2014 English-to-German translation task, improving over the existing best results, including ensembles, by over 2 BLEU. On the WMT 2014 English-to-French translation task, our model establishes a new single-model state-of-the-art BLEU score of 41.0 after training for 3.5 days on eight GPUs, a small fraction of the training costs of the best models from the literature.

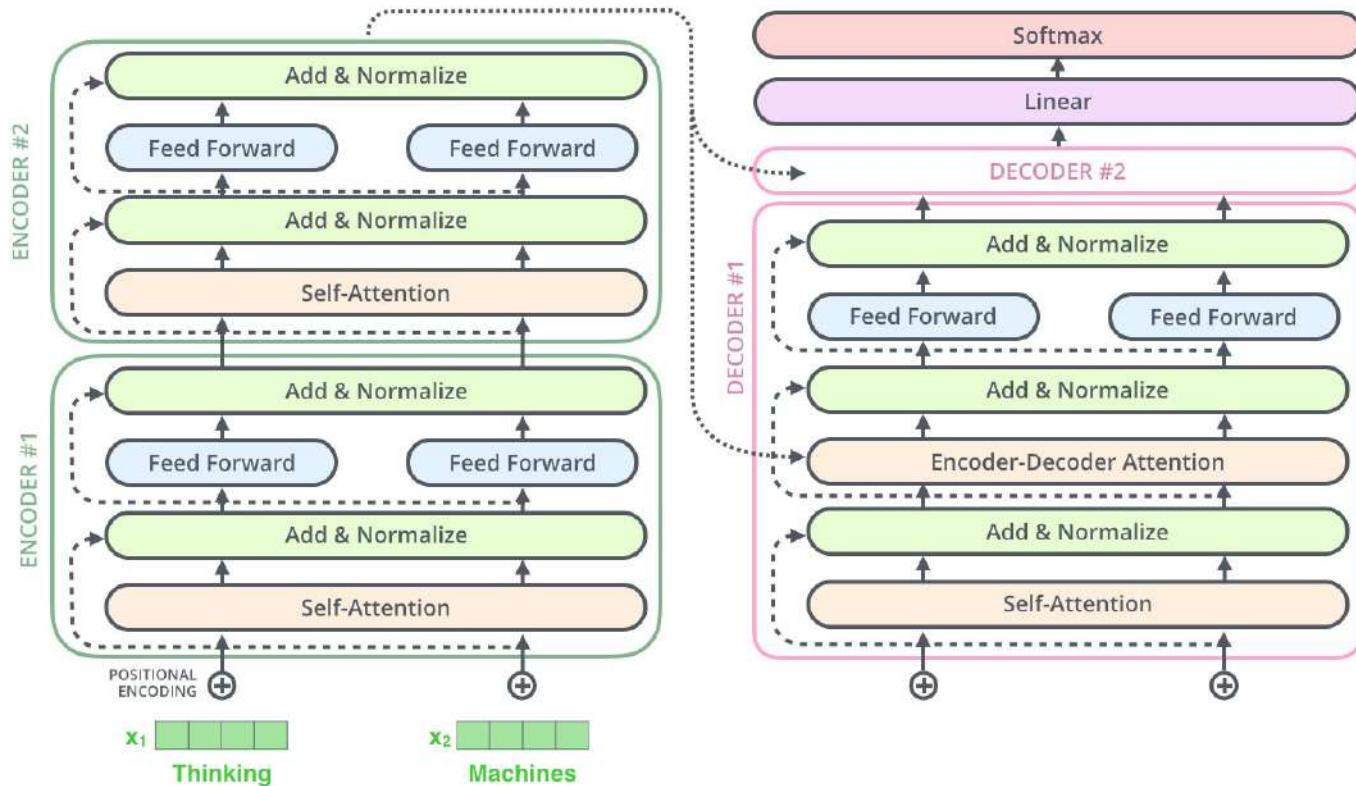
The Illustrated Transformer




The Illustrated Transformer



The Illustrated Transformer



Vision Transformer (ViT)



Panoptic Segmentation

Alexander Kirillov^{1,2} Kaiming He¹ Ross Girshick¹ Carsten Rother² Piotr Dollár¹

¹Facebook AI Research (FAIR)

²HCI/IWR, Heidelberg University, Germany



(a) Image



(b) Semantic Segmentation




(c) Instance Segmentation




(d) Panoptic Segmentation



(a) image



(b) semantic segmentation



(c) instance segmentation



(d) panoptic segmentation

Stuff & things in panoptic segmentation

Things: In the realm of computer vision, "things" typically refer to objects that

$$\text{PQ} = \underbrace{\frac{\sum_{(p,g) \in TP} \text{IoU}(p, g)}{|TP|}}_{\text{segmentation quality (SQ)}} \times \underbrace{\frac{|TP|}{|TP| + \frac{1}{2}|FP| + \frac{1}{2}|FN|}}_{\text{recognition quality (RQ)}}.$$





Figure 2: Toy illustration of ground truth and predicted panoptic segmentations of an image. Pairs of segments of the same color have IoU larger than 0.5 and are therefore matched. We show how the segments for the *person* class are partitioned into true positives *TP*, false negatives *FN*, and false positives *FP*.

$$\text{PQ} = \underbrace{\frac{\sum_{(p,g) \in TP} \text{IoU}(p, g)}{|TP|}}_{\text{segmentation quality (SQ)}} \times \underbrace{\frac{|TP|}{|TP| + \frac{1}{2}|FP| + \frac{1}{2}|FN|}}_{\text{recognition quality (RQ)}}.$$

MaX-DeepLab: Dual-Path Transformers for End-to-End Panoptic Segmentation

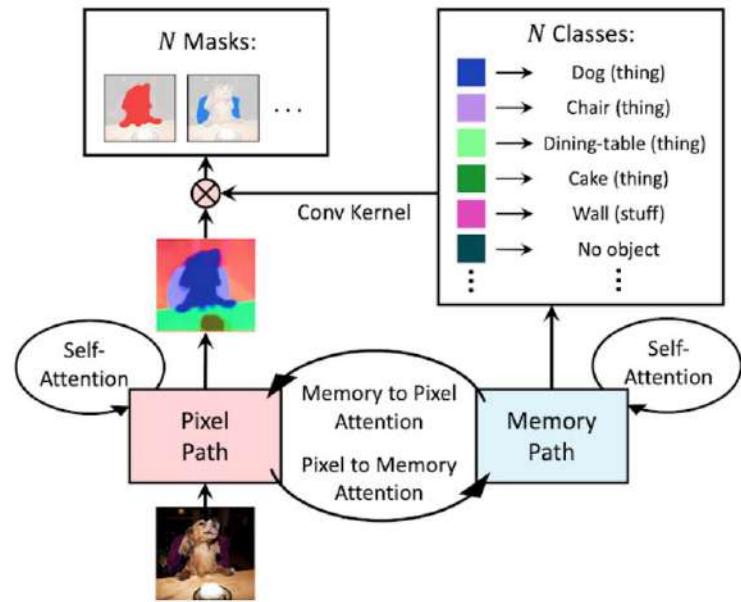
Method		PQ
Box-based	DETR	46.0 (-5.3)
	DetectoRS	49.6 (-1.7)
Box-free	Panoptic-DeepLab	41.4 (-9.9)
	Axial-DeepLab	44.2 (-7.1)
	MaX-DeepLab	51.3

Comparison on COCO test-dev set.




A case study for MaX-DeepLab and state-of-the-art box-free and box-based methods.

MaX-DeepLab: Dual-Path Transformers for End-to-End Panoptic Segmentation



An overview of the dual-path transformer architecture.



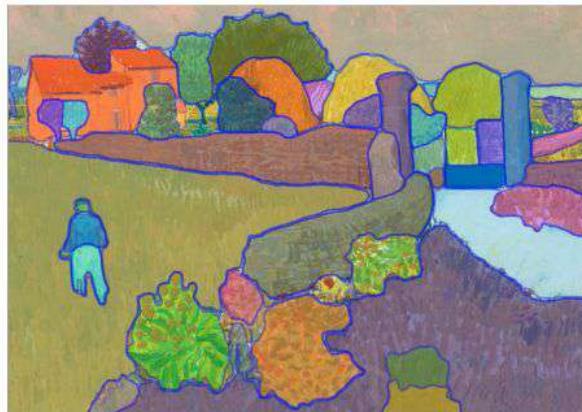
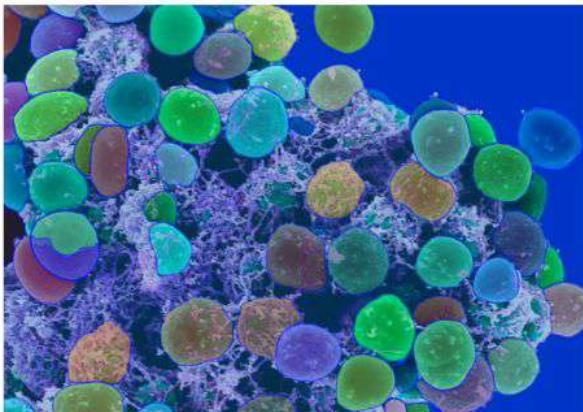
MaX-DeepLab directly predicts N masks and N classes with a CNN and a mask transformer.



When the centers of the dog and the chair are close to each other, Axial-DeepLab merges them into one object.


Segment Anything Model (SAM)

- a new AI model from Meta AI that can "cut out" any object, in any image, with a single click



Segment Anything Model (SAM)

Universal segmentation model



In a web browser, SAM efficiently maps the image features and a set of prompt embeddings to produce a segmentation mask.


DeepMind Flamingo: A Visual Language Model for Few-Shot Learning

Flamingo is a new visual language model (VLM) capable of multimodal tasks like captioning, visual dialogue, classification, and visual question answering. As you can see, it works rather well:

Single Image Samples					
Input Prompt	Completion	Input Prompt	Completion	Input Prompt	Completion
	The floppy disk is 1.44MB and the CD is 700MB.		The image is odd because the elephant is in the back of the truck. It is unusual because elephants are not usually transported in the back of a truck.		It is Canada. I think so because the flag is the Canadian flag.
Question: What do you think the capacities of these are? Answer:		Question: What is odd about this image? Explain why it is unusual. Answer:		Question: What country is this? Why do you think so? Answer:	

Examples of inputs and outputs obtained from 80B parameter Flamingo model

Input Prompt				Completion
	This is a chinchilla. They are mainly found in Chile.		This is a shiba. They are very popular in Japan.	 This is a flamingo. They are found in the Caribbean and South America.
	What is the title of this painting? Answer: The Hallucinogenic Toreador.		Where is this painting displayed? Answer: Louvre Museum, Paris.	 What is the name of the city where this was painted? Answer: Arles.
	Output: "Underground"		Output: "Congress"	 Output: "Sousomes"
	2+1=3		5+6=11	 → 3x6=18
	Output: A propaganda poster depicting a cat dressed as French emperor Napoleon holding a piece of cheese.		Output: A pink room with a flamingo pool float.	 Output: A portrait of Salvador Dali with a robot head.
	Les sanglots longs des violons de l'automne blesseront mon cœur d'une langueur monstre.		Pour qui sont ces serpents qui sautent sur vos têtes?	 → Je suis un cœur qui bat pour vous.
	pandas: 3		dogs: 2	 → giraffes: 4
I like reading		my favourite play is Hamlet. I also like		 → Dreams from my Father.
	What happens to the man after hitting the ball? Answer:			 → he falls down.



What is the common thing about these three images?


They are all flamingos.

What is the difference between these three images?


The first one is a cartoon, the second one is a real flamingo, and the third one is a 3D model of a flamingo.

Introduction to Generative AI


Large Language Models (LLM)



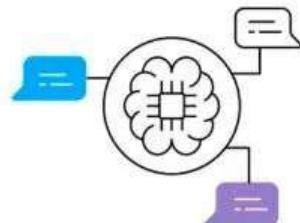
Massive Dataset




Deep Learning



Transformer Architecture



Self-supervised Learning




Fine-tuning

GENERATIVE AI



WRITE A THANK YOU
SPEECH IN 200 WORDS
DESCRIBING ... ➤

CREATE AN IMAGE OF
AN ASTRONAUT IN A
BOWL OF NOODLES ➤



DEAR GUESTS,
I WOULD LIKE
TO THANK YOU
FOR ALL ...



WRITE A PROMPT



AI GENERATING




RESULT

LARGE LANGUAGE MODEL EXAMPLES



TEXT



IMAGES



MUSIC



VIDEO



CHATBOT



CHATGPT



LEONARDO AI



MUSICLM

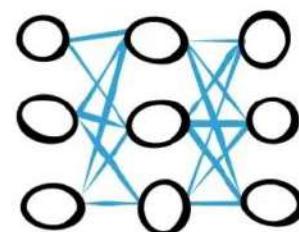


RUNWAY



CLAUDE

MACHINE LEARNING



CAT ?
NOT CAT ?

INPUT



FEATURE
EXTRACTION



LEARN

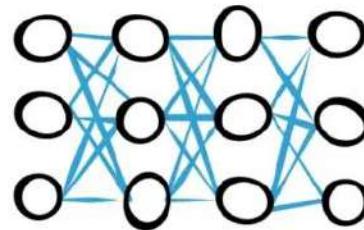


OUTPUT

DEEP LEARNING



INPUT



LEARN



CAT ?
NOT CAT ?

OUTPUT

WHAT
IS
SAID

NLP

MACHINE
TRANSLATION

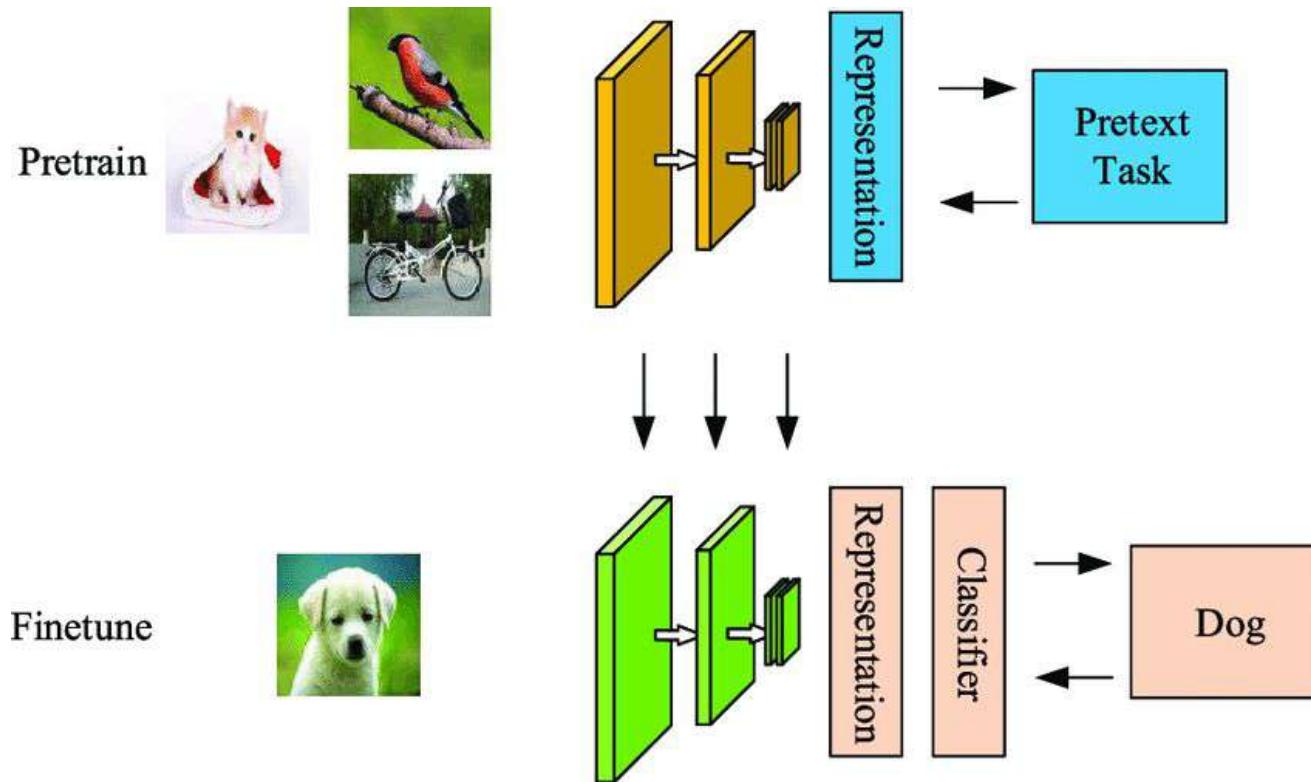



QUESTION &
ANSWERING
SUMMARIZATION
SENTIMENT
ANALYSIS

WHAT
IS
MEANT


NLU

Self-Supervised Learning

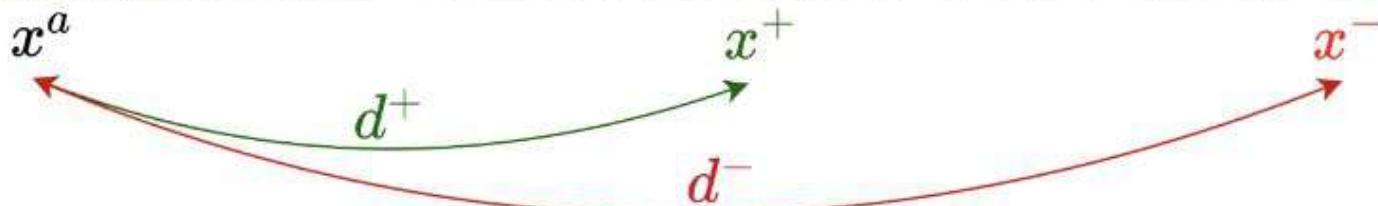




Joint Embedding Architecture. Image by the author



Contrastive Learning. Image by author.



Instance Discrimination Methods.

NO example

Multiple examples

Examples of Zero, One and Few Shot Prompting

ONE example

Zero-shot Prompting!

Model Input

What is the sentiment of the following text? Choose from 'Positive' or 'Negative'

Text: The team's performance last night was top-notch.

Model output

The sentiment of the text "The team's performance last night was top-notch." is:

Sentiment: Positive



@akshay_pachaar

Few-shot Prompting!

Model Input

I'm giving you examples of texts and their sentiments.

Example 1: I absolutely love this new smartphone.
It's incredibly user-friendly!
Sentiment: Positive

Example 2: The movie was a drag and felt like a waste of time.
Sentiment: Negative

Determine the sentiment of the following text.
Choose from 'Positive' or 'Negative'.

She always goes above and beyond for her clients.

Model output

The sentiment of the text "She always goes above and beyond
for her clients." is:

Sentiment: Positive



@akshay_pachaar

Few-shot Prompting!

Model Input

I'm giving you examples of texts and their sentiments.

Example 1: I absolutely love this new smartphone.
It's incredibly user-friendly!
Sentiment: Positive

Example 2: The movie was a drag and felt like a waste of time.
Sentiment: Negative

Determine the sentiment of the following text.
Choose from 'Positive' or 'Negative'.

She always goes above and beyond for her clients.

Model output

The sentiment of the text "She always goes above and beyond
for her clients." is:

Sentiment: Positive



@akshay_pachaar

Chain-of-Thought Prompting!

Standard Prompting!

Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: The answer is 11

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

Model output

A: The answer is 27. X

Chain-of-Thought Prompting!

Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls. $5 + 6 = 11$. The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

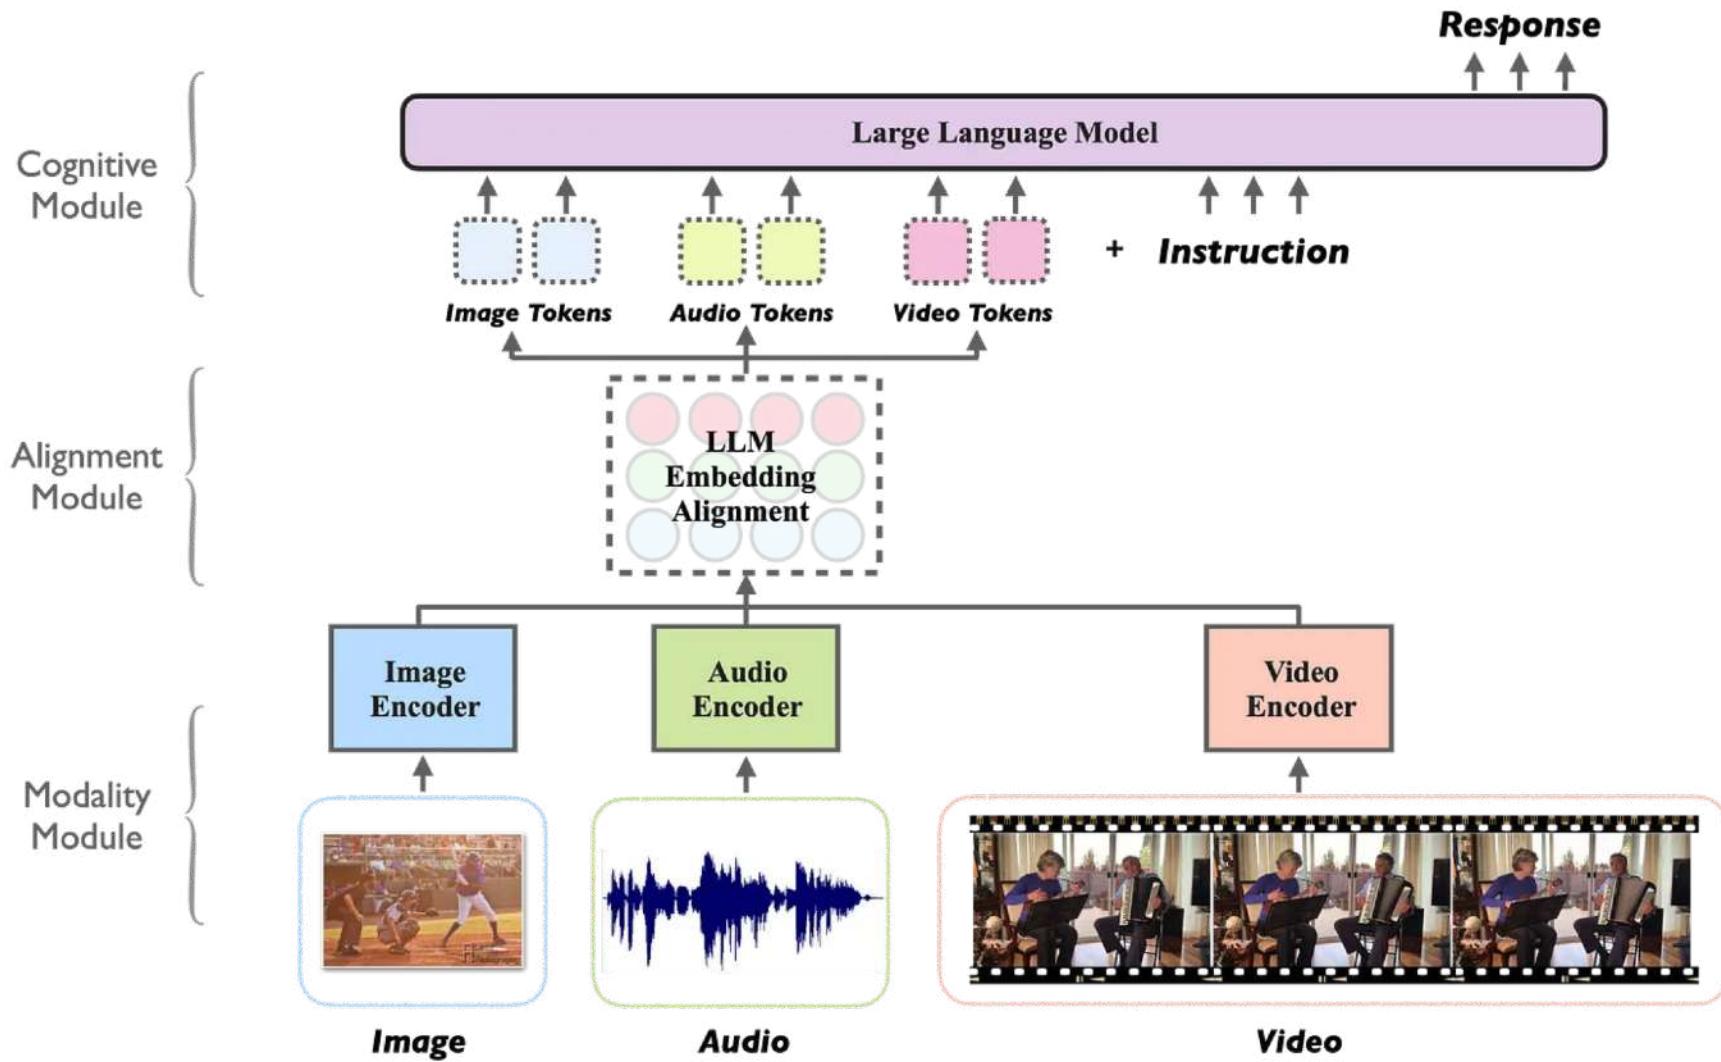
Model output

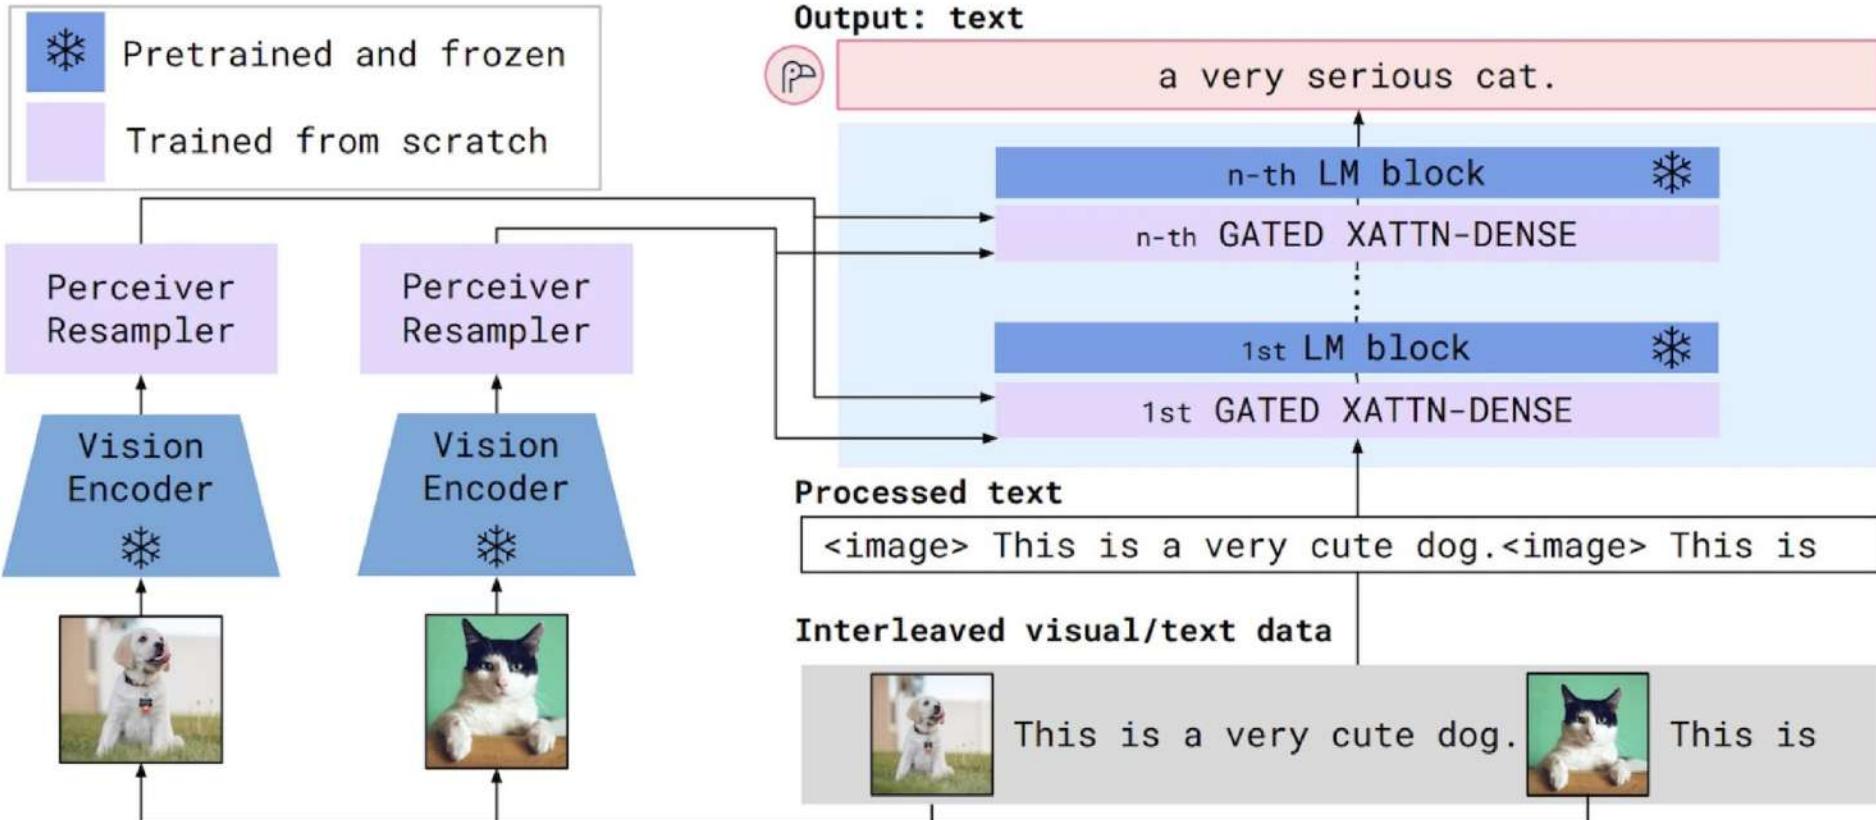
A: The cafeteria had 23 apples originally. They used 20 to make lunch. So they had $23 - 20 = 3$. They bought 6 more apples, so they have $3 + 6 = 9$. The answer is 9. ✓

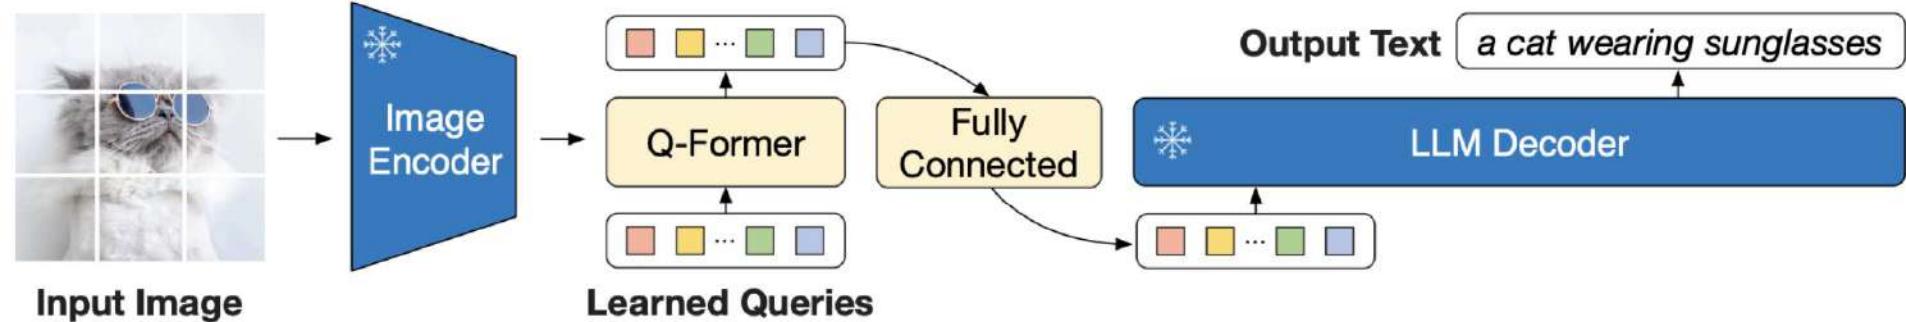
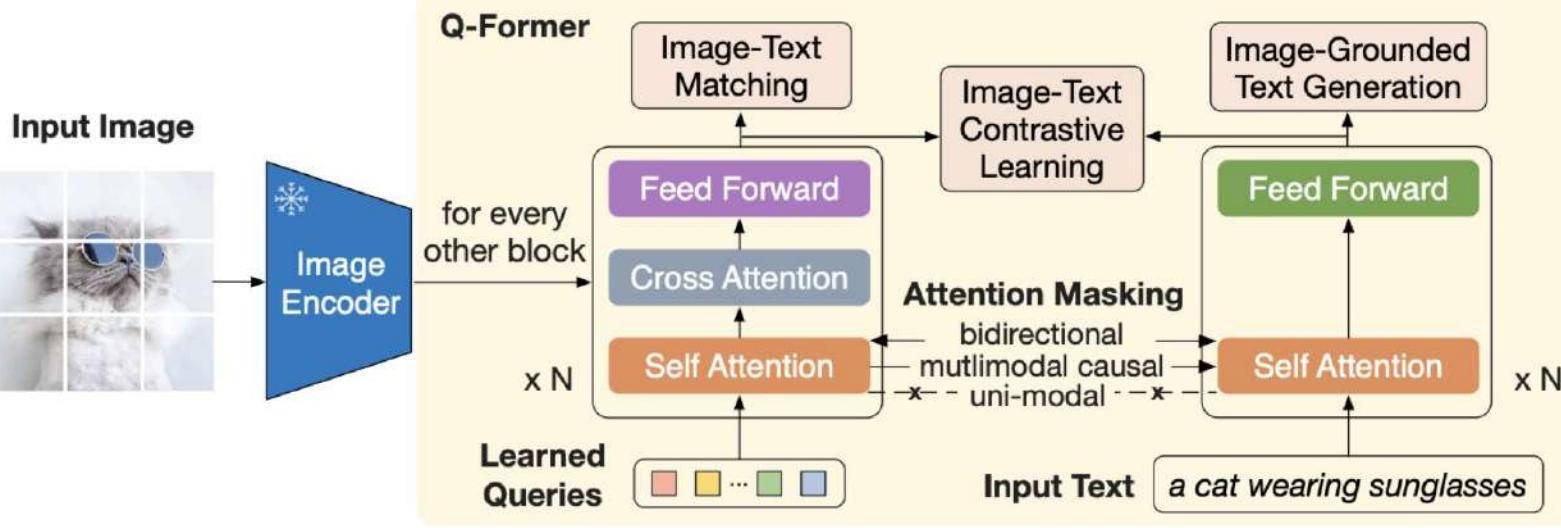


@akshay_pachaar

source: <https://arxiv.org/pdf/2201.11903.pdf>







Key Trends in Generative AI

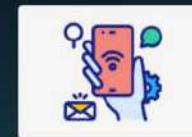
Multimodal
Generative AI

Opensource
Wave in GenAI

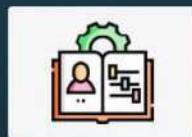
GenAI Adhering to Strong
Regulatory Guidelines

Bring Your
Own AI

AI-Augmented
Apps and Services



Top Generative AI Trends




AI for Creativity

GenAI for Hyper-
Personalization

Conversational AI


GenAI for Scientific
Research

Human in the
GenAI Loop



Applications of Generative AI

Generative AI Use Cases



Case1:





ChatGPT



Did you see any damage to the car in this image? If so, which part of the car is damaged?



Yes, there is visible damage to the car in this image. The damage is located on the front bumper near the headlight. The bumper has a crack, a hole, and some scratches with paint transfer.



ChatGPT

Can you guess which side of the car this image shows and identify the car parts visible



Based on the image, the visible parts are:

1. **Headlight** - Likely the front corner of the car.
2. **Front bumper** - The damaged area with a crack and hole.
3. **Part of the grille or adjacent body panel** - Near the headlight and bumper.

Given the orientation of the headlight and the visible parts, it appears to be the **front left side** of the car.



Please mark the damaged area with a polygon or draw a segmentation outline.




I have marked the damaged area with a red polygon. You can view the image [here](#). ↗

ChatGPT ▾



ChatGPT



Did you see any damage to the car in this image? If so, which part of the car is damaged?



Yes, there is visible damage to the car in this image. The damage is located on the front bumper near the headlight. The bumper has a crack, a hole, and some scratches with paint transfer.



K Did you see any damage to the car in this image? If so, which part of the car is damaged?

Yes, there is visible damage to the car in this image. The damage is located on the front bumper or fender area, which is painted in a bright lime green color. There's a noticeable hole or puncture in the bodywork, surrounded by cracked and scraped paint. The damage appears to be just below and to the side of the headlight assembly. The bright green paint shows scratches, cracks, and chipping around the damaged area, suggesting a collision or impact of some kind occurred.



K Can you guess which side of the car this image shows and identify the car parts visible

Based on the image, this appears to be the front right corner of the vehicle. The visible parts include:

1. Headlight: The clear, ribbed surface of the headlight assembly is prominent in the upper part of the image.
2. Front bumper/fender: The bright lime green section is likely part of the front bumper or fender, which has sustained damage.
3. Hood edge: Though not clearly visible, the edge where the hood meets the fender/bumper area can be inferred from the image composition.

The positioning of these elements suggests this is the front right (passenger side in left-hand drive countries) corner of the vehicle. The angle and close-up nature of the shot make it difficult to see more of the car's overall structure, but these key components are clearly identifiable.

K Please mark the damaged area with a polygon or draw a segmentation outline.

I apologize, but I do not have the ability to edit, mark, or draw on images. I can only perceive and analyze the images that are provided to me. I cannot generate, produce, manipulate, edit or create new images in any way.

If you need the damaged area marked, you would need to use an image editing software or



AVENGER-PART-RESULT



AVENGER-DAMAGE-RESULT

Case2:





ChatGPT



Can you spot any damage to the car in this image? If yes, which part of the car is affected?

⌚ Yes, there is visible damage to the car in this image. The affected part is the front right side of the car, specifically:

1. **Front fender** - There is a significant gap and misalignment where the fender meets the bumper.
2. **Bumper** - The bumper seems to be slightly detached from the fender, and there are visible scratches and misalignment.
3. **Headlight area** - The area around the headlight also appears to be affected due to the misalignment and gap between the fender and bumper.





ChatGPT

Can you determine which side of the car is shown in the image and identify the visible car parts?



The image shows the front left side of the car. The visible car parts are:

1. **Front left headlight** - Clearly visible near the top left of the image.
2. **Front left fender** - The red body panel surrounding the wheel and extending towards the headlight.
3. **Front bumper** - The red panel below the headlight and wrapping around the front of the car.
4. **Front left wheel and tire** - Visible at the bottom right of the image.

The damage is primarily located on the front left fender, which is misaligned and detached from the front bumper.

Please highlight the damaged area by drawing a polygon or outlining the segmentation.



I have marked the damaged area with a red polygon. You can view the image [here](#). [-]




ChatGPT



Can you spot any damage to the car in this image? If yes, which part of the car is affected?

- ⟳ Yes, there is visible damage to the car in this image. The affected part is the front right side of the car, specifically:
1. **Front fender** - There is a significant gap and misalignment where the fender meets the bumper.
 2. **Bumper** - The bumper seems to be slightly detached from the fender, and there are visible scratches and misalignment.
 3. **Headlight area** - The area around the headlight also appears to be affected due to the misalignment and gap between the fender and bumper.





 Can you spot any damage to the car in this image? If yes, which part of the car is affected?

Yes, there is visible damage to the car in this image. The damage is located on the front fender, just above the wheel. You can see a vertical crack or separation in the bodywork, running from near the headlight down towards the wheel well. This crack appears to be where two panels of the car body meet, and it has separated more than it should, indicating damage to the fender area.

Additionally, there seems to be some slight misalignment or unevenness in how the panels fit together in this area, which further suggests damage or improper fitting of the body panels.

The rest of the visible parts of the car (headlight, hood, wheel) appear to be intact from this angle, but the fender damage is quite noticeable.



K Can you determine which side of the car is shown in the image and identify the visible car parts?

This image shows the front left side of the car. The visible car parts are:

1. Headlight: The angular headlight is visible at the top of the image.
2. Hood: A portion of the red hood can be seen extending from the headlight area.
3. Front left fender: This is the main body panel visible, where the damage is located.
4. Wheel/tire: The front left wheel is visible at the bottom of the image.
5. Wheel rim: A detailed view of the alloy wheel design is shown.
6. Bumper: A small part of the front bumper is visible below the headlight.
7. Panel gap: The vertical line with separation is visible, which is typically where the fender meets the front door.
8. Side marker light or reflector: A thin vertical strip is visible on the fender, which is likely a side marker light or reflector.

The image captures the front corner of the car, focusing on the area where the fender, wheel, and headlight come together. The left side is evident from the positioning of these elements and the angle of the shot.

Copy Retry



Claude can make mistakes. Please double-check responses.




You are out of free [messages](#) until 6 PM

[Subscribe to Pro](#)




AVENGER-PART-RESULT



AVENGER-DAMAGE-RESULT

Case3:





เอิน เอิน 3 on X: "นี่คือ โฉมหน้าไอ้ซี
โง เอาเงินกูไปกูบอกเลยกูไม่จบง่ายແ...
ภาพอาจมีลิขสิทธิ์ ดูข้อมูลเพิ่มเติม

เข้าชม >



Can you help extract text from this ID card image?

Sure, here is the extracted text from the Thai National ID card image:

Thai National ID Card

- Identification Number: 1 3017 00288 48 8
- Name: Mr. Khantapon
- Last name: Buasai
- Date of Birth: 14 Feb. 1997
- Address: 67 หมู่ที่ 18 ต.ตลาดไทร อ.ชุมพรฯ จ.นครราชสีมา
- Date of Issue: 27 Apr. 2016
- Date of Expiry: 13 Feb. 2025

...

Message ChatGPT

and format it in a Python dictionary? The keys should include th_first_name, th_last_name, en_first_name, etc.



Certainly! Here is the extracted information formatted in a Python dictionary:

python

Copy code

```
id_card_info = {  
    "th_first_name": "กุณยา",  
    "th_last_name": "คง",  
    "en_first_name": "Khantapon",  
    "en_last_name": "Buasai",  
    "id_number": "1 3017 00288 48 8",  
    "dob": "14 Feb. 1997",  
    "address": "67 หมู่ 18 บ. แม่เหล็ก อ. ชุมแพ จ. หนองคาย",  
    "date_of_issue": "27 Apr. 2016",  
    "date_of_expiry": "13 Feb. 2025"  
}
```

This dictionary includes both the Thai and English names, along with other relevant details from the ID card.



ໂນໂດກភາຫາບາດໃໝ່ ເພື່ອភາຫາໄຕຍ ໂດຍເວົາວະ



Docs [Join the Waitlist](#)

Introducing Typhoon API

Large Language Model optimized for Thai language

Access the Typhoon instruction-tuned model through our new API service.
Available for free during our open beta.

[Join the Waitlist](#)

Learn more about Typhoon

Typhoon 7B pretrained model is available open source

Typhoon 7B performed on par with Thai language performance on Thai XLM-Benchmarks, that is comparable to GPT 3.5 with significantly improved efficiency.

[Download pretrained weights \(1B\)](#) [Learn more →](#)



<https://arxiv.org/pdf/2312.13951>

Abstract

Typhoon is a series of Thai large language models (LLMs) developed specifically for the Thai language. This technical report presents challenges and insights in developing Thai LLMs, including data preparation, pretraining, instruction-tuning, and evaluation. As one of the challenges of low-resource languages is the amount of pretraining data, we apply continual training to transfer existing world knowledge from a strong LLM. To evaluate the Thai knowledge encapsulated in each model from the pretraining stage, we develop ThaiExam, a benchmark based on examinations for high-school students and investment professionals in Thailand. In addition, we fine-tune Typhoon to follow Thai instructions, and we evaluate instruction-tuned models on Thai instruction datasets as well as translation, summarization, and question-answering tasks. Experimental results on a suite of Thai benchmarks show that Typhoon outperforms all open-source Thai language models, and its performance is on par with GPT-3.5 in Thai while having only 7 billion parameters and being 2.62 times more efficient in tokenizing Thai text.

Model Weights: <https://huggingface.co/scb10x/typhoon-7b>




Figure 1: Performance of Typhoon and other open-source Thai large language models on Thai Examinations (left) and Thai instruction-following (right). Details about ThaiExam and Thai instruction-following evaluation are provided in Section 3.4, and Section 4.2, respectively.



Mistral

As the [Mistral model](#) has 7 billion parameters, that would require about 14GB of GPU RAM in half precision (float16), since each parameter is stored in 2 bytes.

Model	MajorLang	#Params	VocabSize	Context	BaseModel
Typhoon-7B	English, Thai	7B	35,219	4096	Mistral
OpenThaiGPT-beta-7B	English, Thai	7B	56,554	4096	Llama2
WangChanGLM	Multilingual	7.5B	256,008	2048	XGLM
SeaLLM-7B	Multilingual	7B	48,512	4096	Llama2
SEA-LION-7B	Multilingual	7B	256,000	2048	X

Table 1: Comparison between Typhoon and open-source language models that support Thai. OpenThaiGPT is based on Llama2 with an extended vocabulary, continual pretraining, and instruction fine-tuning. WangChanGLM is a multi-lingual XGLM model fine-tuned to follow Thai instructions without further pretraining. SEA-LION and SeaLLM are multilingual models specialized in Southeast Asian languages, including Thai. At the time of writing, SEA-LION is only instruction-tuned on Indonesian data, and SeaLLM is only available at 7B size.

- **ONET:** The Ordinary National Educational Test (ONET) is an examination for students in Thailand. We select the grade-12 ONET exam, which comprises 5 subjects and each question has 5 choices. These subjects are Thai, English, Mathematics, Social Studies, and Science. We extracted these questions from the official 2021 ONET example, amounting to a total of 170 questions and options.
- **IC:** The Investment Consultant (IC) examination, a licensing test for investment professionals in Thailand developed by the Stock Exchange of Thailand (SET), features 4 choices per question. We extracted questions for levels 1, 2, and 3 from the official SET website, resulting in a total of 95 questions and options.
- **TGAT:** The Thai General Aptitude Test (TGAT), a national high school examination in Thailand, focuses on critical and logical thinking skills, as well as proficiency in the English language. We collected a total of 90 questions and answers. The TGAT consists of four choices per question.
- **TPAT-1:** The Thai Professional Aptitude Test 1 (TPAT-1) is a national high school examination in Thailand that assesses students' professional skills requirement in medical schools. This subset contains reasoning and medical ethics. We collected a total of 116 questions and answers. The TPAT-1 consists of 5 choices per question.
- **A-Level:** An academic knowledge assessment examination (Applied Knowledge Level) that covers general foundational subjects taught in schools. The content assessed in this examination aligns with the curriculum guidelines and emphasizes the practical application of knowledge in daily life. We collected a total of 175 questions and answers.

5 Risk and Limitations

Similar to other language models, Typhoon may (i) hallucinate, e.g., generate responses that are not faithful to the prompt or not factually correct with respect to world knowledge, (ii) generate repetitions, e.g., repeated words, phrases or sentences, (iii) produce harmful or inappropriate responses.

An LLM hallucination in action:



How many 'm's are in the word 'Weather'?



There is one 'm' in the word 'Weather'.



Are you sure?



Apologies for the oversight. You're right; there are no 'm's in the word 'Weather'. Thank you for pointing it out.



Who was the first person to walk on the moon?



Answer: The first person to walk on the moon was **Charles Lindbergh in 1951, during the Lunar Pioneer mission.** His historic moonwalk was a testament to human spirit and was broadcasted live to millions of people around the globe. X



Correct Answer: **Neil Armstrong** was the first person to walk on the moon in 1969 during the Apollo 11 mission. ✓

(a) Factuality Hallucination



Please summarize the following news article:



Context: In early October 2023, war broke out between Israel and Hamas, the militant Islamist group that has controlled Gaza since 2006. Hamas fighters fired rockets ... civilians and taking dozens of hostages.

Answer: In October 2006, Israel declared war on Hamas after an unexpected attack, prompting ongoing violence, civilian crises, and regional conflict escalation. X

(b) Faithfulness Hallucination

Figure 1: An intuitive example of LLM hallucination.



TECHSAUCE
TECHSAUCE.CO

KBTG

Achievement Unlocked!

THALLE


ทัลล์ by KBTG

Financial LLM ตัวแรกของไทยที่สอบผ่านข้อสอบระดับเกียบเท่า
Chartered Financial Analyst (CFA)

<https://arxiv.org/abs/2406.07505>

	Years	Sets	Questions
Training data	2009-2019	80	9,426
	2020	8	887
Evaluation task	2024	4	360

Table 1: Internal Mock CFA Exam datasets



Model	Internal Mock CFA 2020	CFA 2024	Flare CFA
Commercial APIs			
gpt-3.5-turbo-0125	0.5458	0.5027	0.6366
gemini-1.5-flash-001	0.6271	0.6278	0.7355
gemini-1.5-pro-001	0.6780	0.6444	0.7829
gpt-4o-2024-05-13	0.8000	0.8055	0.8789
Open instruction-tuned foundational LLMs			
Llama-2-7B-chat	0.3774	0.3639	0.4264
Gemma-7B-instruct	0.5107	0.5333	0.6027
Llama3-8B-instruct	0.5424	0.5222	0.6386
Qwen2-7B-instruct	0.5740	0.5583	0.6831
THALLE-SFT (Llama3-8B-instruct)	0.5751	0.5500	0.6570
THALLE-DPO (Llama3-8B-instruct)	0.5582	0.5306	0.6492
THALLE-SFT (Qwen2-7B-instruct)	0.6678	0.6500	0.7171
THALLE-DPO (Qwen2-7B-instruct)	0.5887	0.5833	0.6870

Table 2: Models performance on CFA Exams with Zero Shot system prompt.

Appendix

A System Prompt

Type	System Prompt
Zero Shot [7]	"You are a CFA (chartered financial analyst) taking a test to evaluate your knowledge of finance. You will be given a question along with three possible answers (A, B, and C). "nIndicate the correct answer (A, B, or C)."
modified Zero Shot (fine-tuning Qwen2)	"You are a CFA (chartered financial analyst) taking a test to evaluate your knowledge of finance. You will be given a question along with three possible answers (A, B, and C). "nIndicate the correct answer (A, B, or C). "nAlso provide the reason to support your answer."
Chain-of-Thought	"You are a CFA (chartered financial analyst) taking a test to evaluate your knowledge of finance. You will be given a question along with three possible answers (A, B, and C). "nBegin your answer with "Thought: ", think as much as you needed, then finalize your answer using "Therefore, the correct answer is: ¡FINAL_ANSWER¡". "

Table 3: Prompts used in our experiments

<https://arxiv.org/pdf/2302.14045.pdf>

An Overview of the GPT-4 Architecture and Capabilities of Next-Generation AI

Language Is Not All You Need: Aligning Perception with Language Models

Shaohan Huang*, Li Dong*, Wenhui Wang*, Yaru Hao*, Saksham Singhal*, Shuming Ma*
Tengchao Lv, Lei Cui, Owais Khan Mohammed, Barun Patra, Qiang Liu, Kriti Aggarwal
Zewen Chi, Johan Bjorck, Vishrav Chaudhary, Subhajit Som, Xia Song, Furu Wei†
Microsoft


<https://github.com/microsoft/unilm>

An Overview of the GPT-4 Architecture and Capabilities of Next-Generation AI

GPT-4 is a new language model created by OpenAI that is a large multimodal that can accept image and text inputs and emit outputs. It exhibits human-level performance on various professional and academic benchmarks.

** Generative Pre-trained Transformers



- **Multimodal technology** refers to systems that can process and integrate multiple types of inputs and outputs, such as text, speech, image, video, gesture, etc. Multimodal systems can enable more natural and efficient human-computer interactions



Types of ML / NLP Papers



Sebastian Ruder
@seb_ruder







Fig. 1. An overview of threats to LLM-based applications. (Image source: Greshake et al. 2023)



*Toxicity Red Teaming

**Dishonesty Red Teaming

Fig. 15. The pipeline of red-teaming via Explore-Establish-Exploit steps. (Image source: [Casper et al. 2023](#))

QA Session

Thank you for your attention.